Final Project Report: Time Series Analysis

Time Series Forecasting and Modeling: Analyzing NonSeasonal Dataset — A Study on Baltic Dry Index —

Achintya
Student ID: 20021326

April 2025
Course: MA 641 - Time Series Analysis I
Term: Fall 2025

Contents

1 Introduction and Motivation **1**

2 Data Description and Preprocessing **1**
- 2.1 Baltic Dry Index (BDI) Dataset . . . 1
- 2.2 Data Loading and Initial Structure . . . 2
- 2.3 Preprocessing and Time Series Conversion . . . 2

3 Exploratory Data Analysis (EDA) **3**
- 3.1 Line Plot of Monthly BDI Prices . . . 3
- 3.2 Plot with Trend Line (LOESS Smoother) . . . 3

4 Stationarity Testing and Differencing **4**
- 4.1 Augmented Dickey-Fuller Test on Original Series (bdi_ts) . . . 4
- 4.2 First-Order Differencing and ADF Retest . . . 4

5 Model Identification (ACF & PACF) **5**
- 5.1 ACF & PACF Plots of bdi_diff . . . 5
- 5.2 Interpretation . . . 6

6 Model Selection and Fitting **7**
- 6.1 Auto ARIMA Selection using auto.arima() . . . 7
- 6.2 Final Model Choice . . . 7

7 Residual Diagnostics for ARIMA(0,1,5) **8**

8 Forecasting with ARIMA(0,1,5) **9**
- 8.1 ARIMA Forecast (12 Months Ahead) . . . 9

9 Volatility Modeling using GARCH **11**
- 9.1 GARCH Model Setup and Fitting . . . 12
- 9.2 Conditional Volatility Forecast . . . 13

10 Model Validation **14**

11 Model Comparison Summary **15**

12 Interpretations and Recommendations **16**

13 Limitations and Future Work **16**
- 13.1 Limitations . . . 16
- 13.2 Future Work . . . 17

14 Executive Summary **17**

15 Conclusions **18**

16 References **19**

1 Introduction and Motivation

In today's data-driven economy, the ability to model and forecast time-dependent phenomena plays a crucial role across various industries. Time series forecasting empowers decision-makers to detect structure, understand historical behavior, and project future values with improved accuracy.

This project focuses on the Baltic Dry Index (BDI) — a nonseasonal, high-frequency economic indicator used to assess the cost of shipping raw materials like coal, iron ore, and grain across global sea routes. Because it captures international trade volumes and supply-demand dynamics, BDI is considered a barometer of global economic health.

Our goal is to model and forecast the BDI using classical Box-Jenkins ARIMA methodology, supplemented by GARCH modeling to capture the volatility clustering typical of financial time series. This analysis not only enhances predictive accuracy but also adds interpretive depth for economic insights.

Our objectives are:

- Conduct in-depth preprocessing and visualization of the BDI series.
- Assess stationarity and apply transformations as necessary.
- Identify optimal ARIMA model structures using ACF, PACF, and auto.arima.
- Fit and validate selected models using residual diagnostics.
- Forecast BDI values and assess confidence bounds.
- Augment forecasts with GARCH-based volatility insights.
- Interpret findings in the context of economic implications.

2 Data Description and Preprocessing

2.1 Baltic Dry Index (BDI) Dataset

- **Source:** Investing.com BDI Historical Data
- **Period Covered:** January 2002 to January 2024
- **Frequency:** Daily, aggregated to monthly for this analysis.
- **Variables Used:**
  - Date: Observation date
  - Price: BDI closing price
- **Characteristics:** This dataset reflects nonseasonal, high-volatility patterns typical of financial indicators. Aggregation to monthly frequency smooths out short-term noise while preserving structural trends needed for modeling.

2.2 Data Loading and Initial Structure

The necessary R packages (`forecast`, `tseries`, `ggplot2`, `lubridate`, `zoo`, `scales`, `rugarch`) were loaded. The BDI data was read from the CSV file.

```
> str(bdi_data)
'data.frame':   265 obs. of  7 variables:
 $ Date    : chr  "01/01/2024" "12/01/2023" "11/01/2023" "10/01/2023" ...
 $ Price   : chr  "1,398.00" "2,094.00" "2,937.00" "1,459.00" ...
 $ Open    : chr  "2,093.00" "3,192.00" "1,401.00" "1,737.00" ...
 $ High    : chr  "2,093.00" "3,192.00" "1,401.00" "1,737.00" ...
 $ Low     : chr  "2,093.00" "3,192.00" "1,401.00" "1,737.00" ...
 $ Vol.    : logi  NA NA NA NA NA NA ...
 $ Change..: chr  "-33.24%" "-28.70%" "101.30%" "-14.23%" ...
> head(bdi_data)
        Date    Price     Open     High      Low Vol. Change..
1 01/01/2024 1,398.00 2,093.00 2,093.00 2,093.00   NA  -33.24%
2 12/01/2023 2,094.00 3,192.00 3,192.00 3,192.00   NA  -28.70%
3 11/01/2023 2,937.00 1,401.00 1,401.00 1,401.00   NA  101.30%
4 10/01/2023 1,459.00 1,737.00 1,737.00 1,737.00   NA  -14.23%
5 09/01/2023 1,701.00 1,065.00 1,065.00 1,065.00   NA   56.63%
6 08/01/2023 1,086.00 1,150.00 1,150.00 1,150.00   NA   -3.64%
```

Figure 1: Structure of raw ice_cream_data

Note: Expected output shows Date as character, Price with commas, etc.

2.3 Preprocessing and Time Series Conversion

- **Date Conversion:** Dates parsed to Date type (`%m/%d/%Y`).
- **Sorting:** Data sorted chronologically.
- **Numeric Conversion:** Commas removed from Price, converted to numeric.
- **NA Handling:** Rows with NA in Price removed. Only Date and Price kept.
- **Monthly Time Series Creation:** Cleaned data converted to `ts` object (`bdi_ts`) with `frequency=12`, starting Jan 2002.

Table 1: Example of Cleaned BDI Data Head

Index	Date	Production Index
265	2002-01-01	908
264	2002-02-01	1031
263	2002-03-01	1082
262	2002-04-01	1027
261	2002-05-01	1042
260	2002-06-01	1005

Note: Expected output shows Date and numeric Price columns

3 Exploratory Data Analysis (EDA)

Visualizing the time series helps identify patterns like trend, seasonality (or lack thereof), and volatility.

3.1 Line Plot of Monthly BDI Prices

[Line plot of BDI monthly time series]

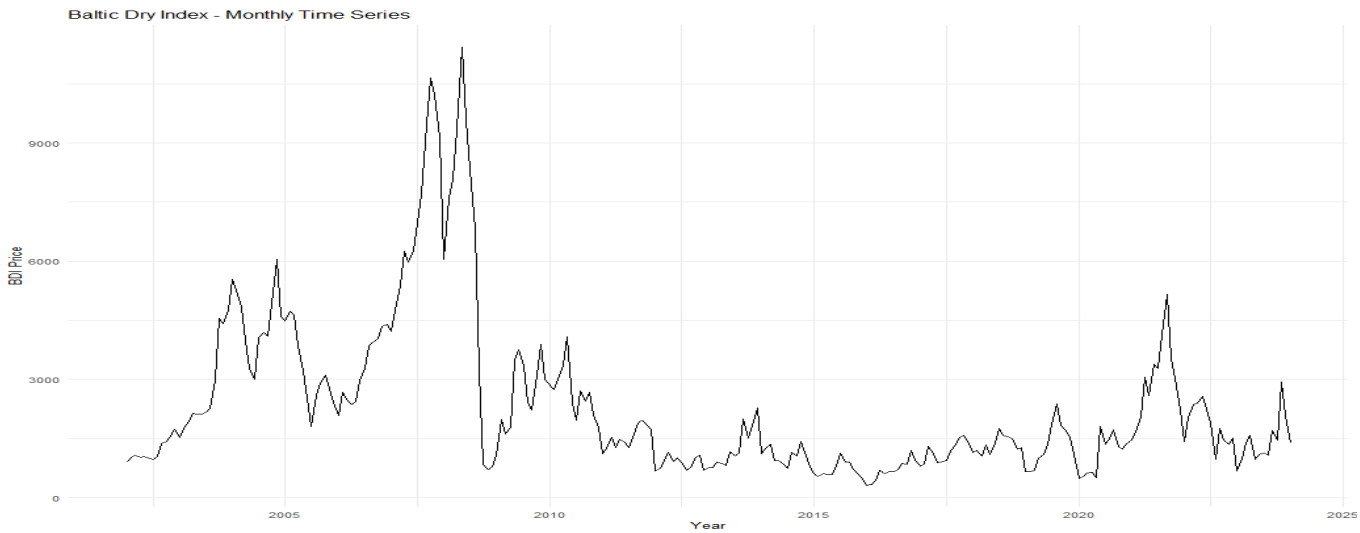

Figure 2: Baltic Dry Index - Monthly Time Series (autoplot)

[Detailed BDI plot (ggplot2)]

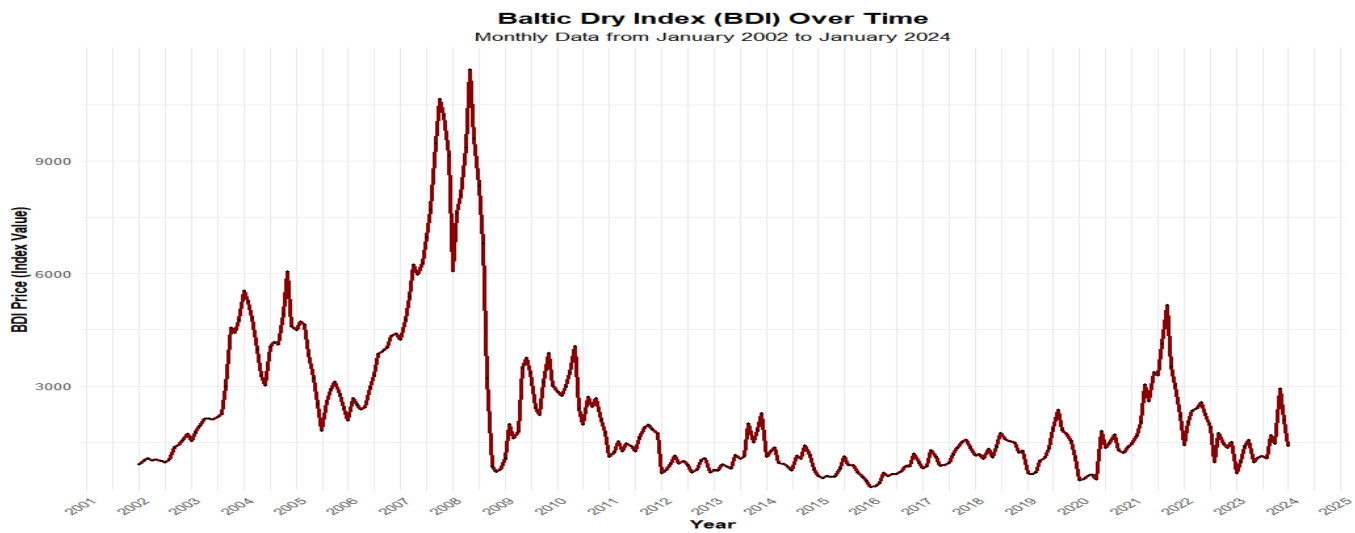

Figure 3: Baltic Dry Index (BDI) Over Time (ggplot2)

3.2 Plot with Trend Line (LOESS Smoother)

[BDI plot with LOESS smoother]

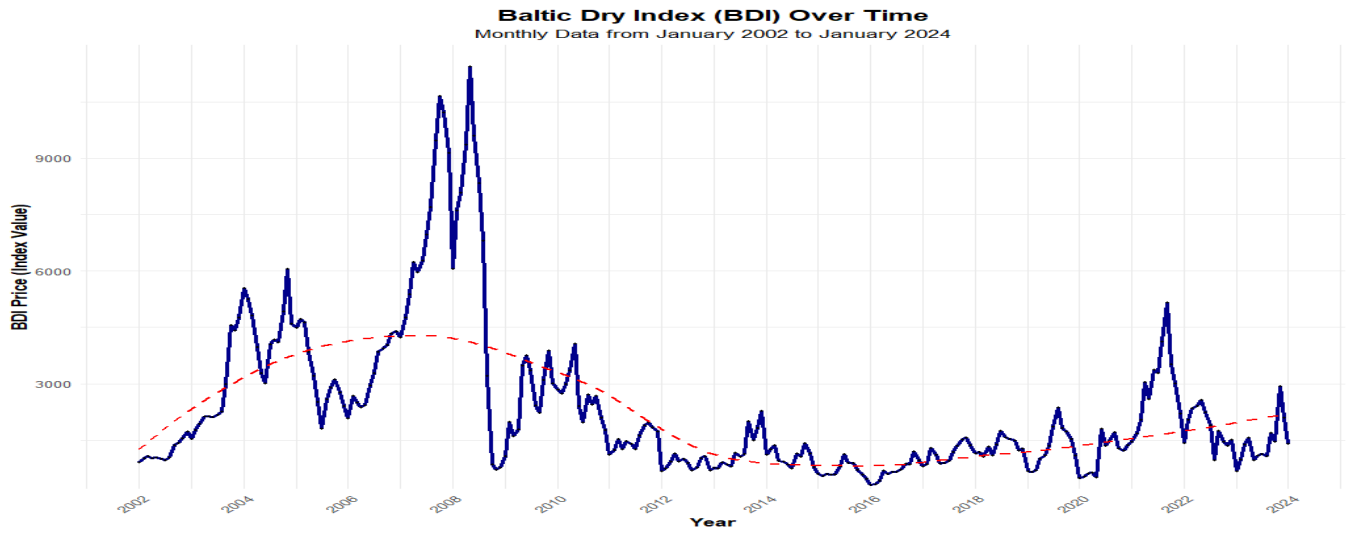

Figure 4: Baltic Dry Index (BDI) Over Time with Trend (LOESS Smoother)

Interpretation: The plots reveal strong trend components and significant volatility. BDI fluctuates sharply, with notable peaks and troughs, particularly around the 2008 financial crisis. No clear repeating seasonal pattern is evident, validating the choice of nonseasonal modeling approaches. The LOESS smoother effectively captures the underlying long-term movements in the index. The visual evidence strongly suggests non-stationarity.

4 Stationarity Testing and Differencing

ARIMA models require stationary data. We use the Augmented Dickey-Fuller (ADF) test to check for unit roots.

4.1 Augmented Dickey-Fuller Test on Original Series (bdi_ts)

ADF Test Result (Original):

```
    Augmented Dickey-Fuller Test

data:  bdi_ts
Dickey-Fuller = -3.249, Lag order = 6, p-value = 0.08013
alternative hypothesis: stationary
```

Interpretation: The p-value (0.08013) is greater than 0.05. We fail to reject the null hypothesis of non-stationarity. The BDI series requires differencing. "This result is expected—most real-world economic time series are not stationary initially due to trends or structural shifts."

4.2 First-Order Differencing and ADF Retest

We apply first-order differencing ($d = 1$) and re-test.

```
# R code for First-Order Differencing and ADF Retest goes here
# e.g., diff(), autoplot(), adf.test()
```

[Plot of first-differenced BDI series (bdi_diff)]

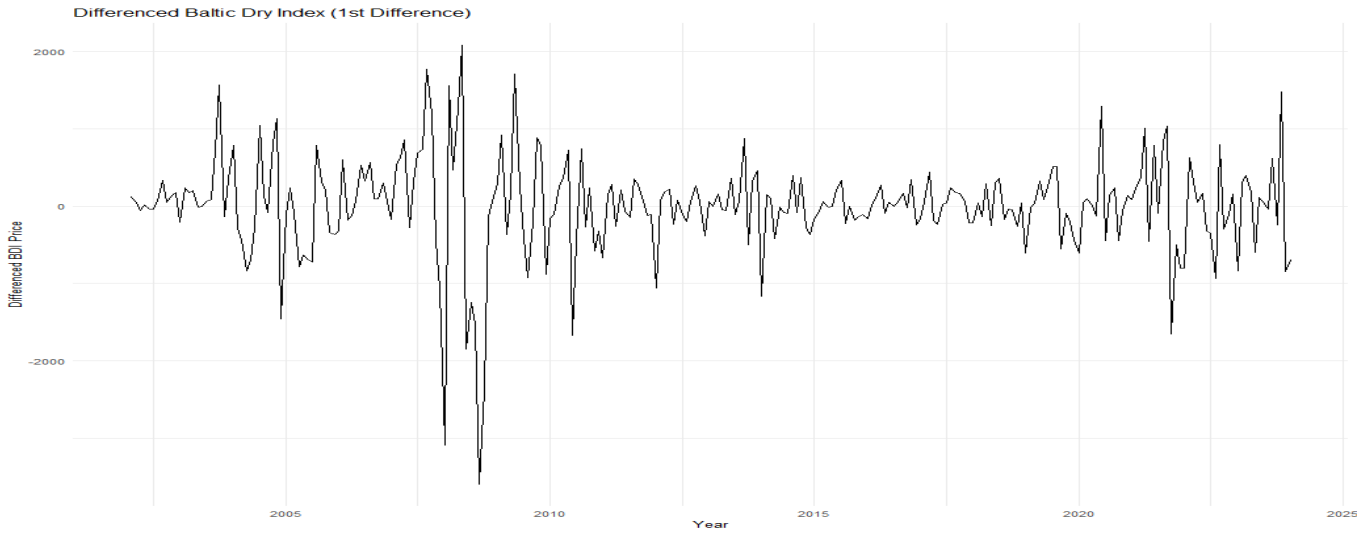

Figure 5: Differenced Baltic Dry Index (1st Difference)

ADF Test Result (Differenced):

```
Augmented Dickey-Fuller Test

data:  bdi_diff
Dickey-Fuller = -6.6623, Lag order = 6, p-value = 0.01
alternative hypothesis: stationary
```

Interpretation: The p-value (0.01) is less than 0.05. We reject the null hypothesis. The first-differenced BDI series (`bdi_diff`) is stationary. Therefore, we set $d = 1$ for our ARIMA models.

5 Model Identification (ACF & PACF)

We examine the ACF and PACF plots of the stationary (differenced) series `bdi_diff` to identify potential AR(p) and MA(q) orders.

5.1 ACF & PACF Plots of bdi_diff

[ACF plot of differenced BDI series]

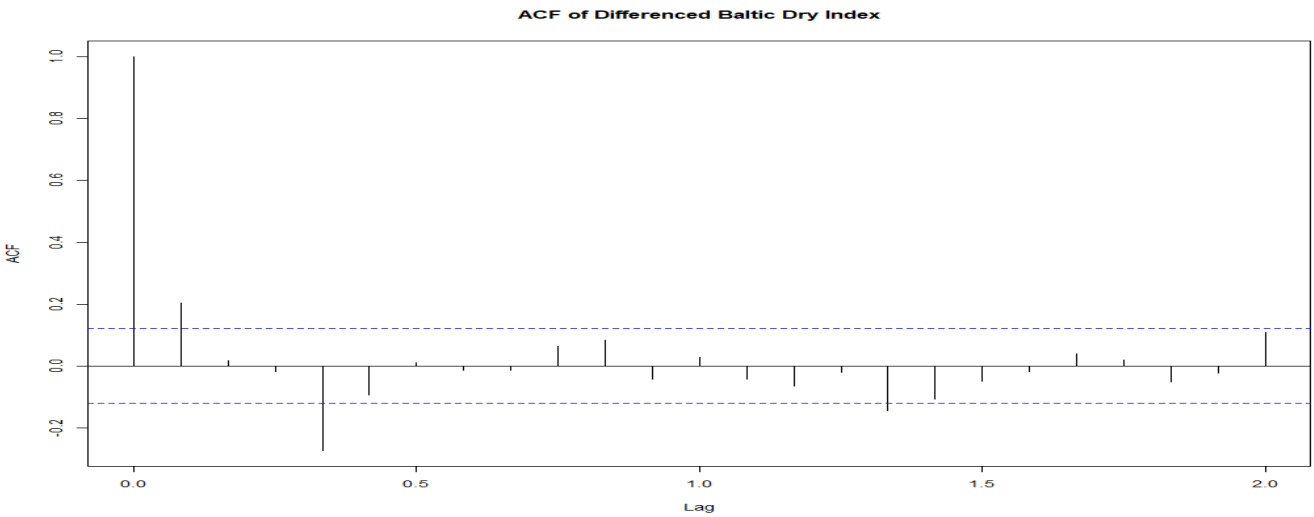

Figure 6: ACF of Differenced Baltic Dry Index

[PACF plot of differenced BDI series]

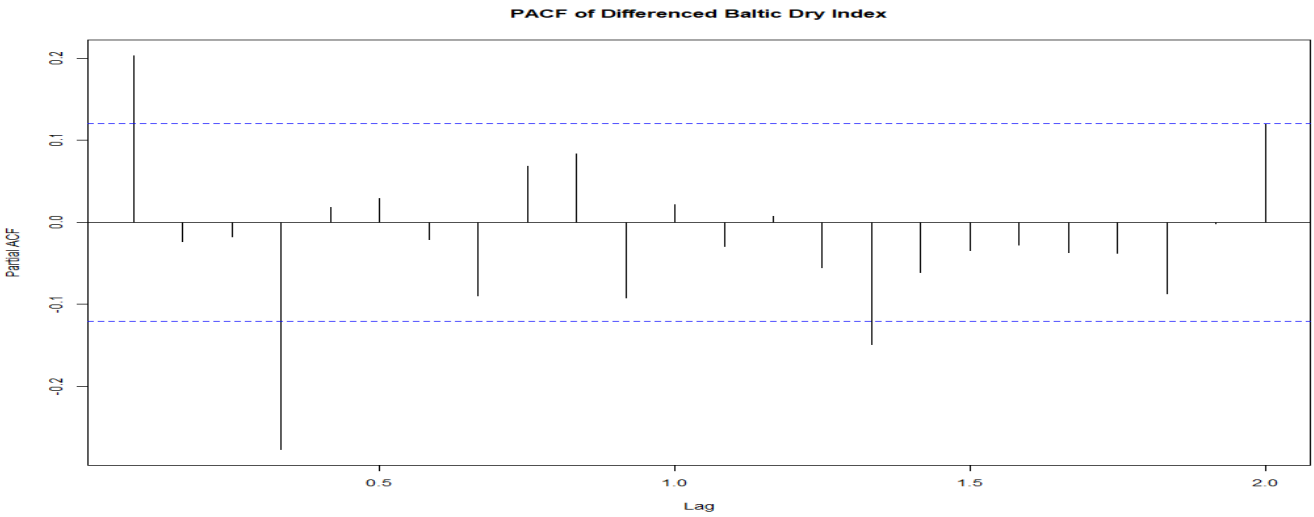

Figure 7: PACF of Differenced Baltic Dry Index

5.2 Interpretation

Based on the visual inspection of the ACF and PACF plots:

Plot Type	Pattern Observed	Interpretation Guide	Suggested Order
ACF	Significant spike at lag 1, cuts off	MA(q) model	q=1?
PACF	Significant spike at lag 1, cuts off	AR(p) model	p=1?

The plots show sharp cutoffs after lag 1 in both the ACF and PACF. This suggests potential models like MA(1), AR(1), or a mixed ARMA(1,1) structure for the differenced series. "Both ACF and PACF support a combined ARIMA model, possibly ARIMA(1,1,1)."

Preliminary Manual Model: Combining this with $d = 1$, a preliminary model based purely on this visual inspection is ARIMA(1,1,1).

6 Model Selection and Fitting

We compare the manually suggested model with automatically selected models using information criteria (AICc).

6.1 Auto ARIMA Selection using auto.arima()

The `auto.arima` function searches for the optimal (p, d, q) order. We set `stepwise=FALSE` and `approximation=FALSE` for a more thorough search.

Output Trace Snippet (showing AICc values):

```
...
ARIMA(0,1,4) : 4135.498
ARIMA(0,1,5) : 4134.635  <- Lowest AICc
ARIMA(1,1,1) : 4152.152
...
Best model: ARIMA(0,1,5)
```

[Full auto.arima summary output table for ARIMA(0,1,5)]

Table 2: ARIMA(0,1,5) Model Summary for `bdi_ts`

Coefficient	**MA1**	**MA2**	**MA3**	**MA4**	**MA5**
Estimate	0.2150	0.0612	0.0174	-0.3289	-0.1041
Std. Error	0.0631	0.0619	0.0583	0.0651	0.0644

Sigma^2	360516
Log Likelihood	-2061.32
AIC	4134.64
AICc	4134.97
BIC	4156.09

Metric	**ME**	**RMSE**	**MAE**	**MPE**	**MAPE**	**MASE**	**ACF1**
Training Set	1.6924	593.59	391.19	-3.912	20.715	0.283	-0.0060

Ljung-Box Test (Residuals)	
Q*	21.983
Degrees of Freedom	19
p-value	0.2851
Model df	5
Total Lags Used	24

Selected Model: `auto.arima` selects ARIMA(0,1,5) as the best model based on the lowest AICc value (4134.64).

6.2 Final Model Choice

Comparing the manually proposed ARIMA(1,1,1) (AICc $\approx$ 4152.15) and the automatically selected ARIMA(0,1,5) (AICc $\approx$ 4134.64), the ARIMA(0,1,5) model provides a significantly better fit according to the information criterion. Furthermore, residual diagnostics for the ARIMA(1,1,1) yielded a Ljung-Box test p-value of 0.03039, suggesting some remaining autocorrelation, whereas the ARIMA(0,1,5) model resulted in a Ljung-Box p-value of 0.2851, indicating residuals closer to white noise. Based on both criteria, we proceed with ARIMA(0,1,5) as the final model for mean forecasting.

7 Residual Diagnostics for ARIMA(0,1,5)

We check if the residuals of the fitted ARIMA(0,1,5) model behave like white noise (uncorrelated, zero mean, constant variance, normally distributed). The ARIMA(0,1,5) model included five MA terms and captured the differenced series dynamics well.

[Residual time plot for ARIMA(0,1,5)]

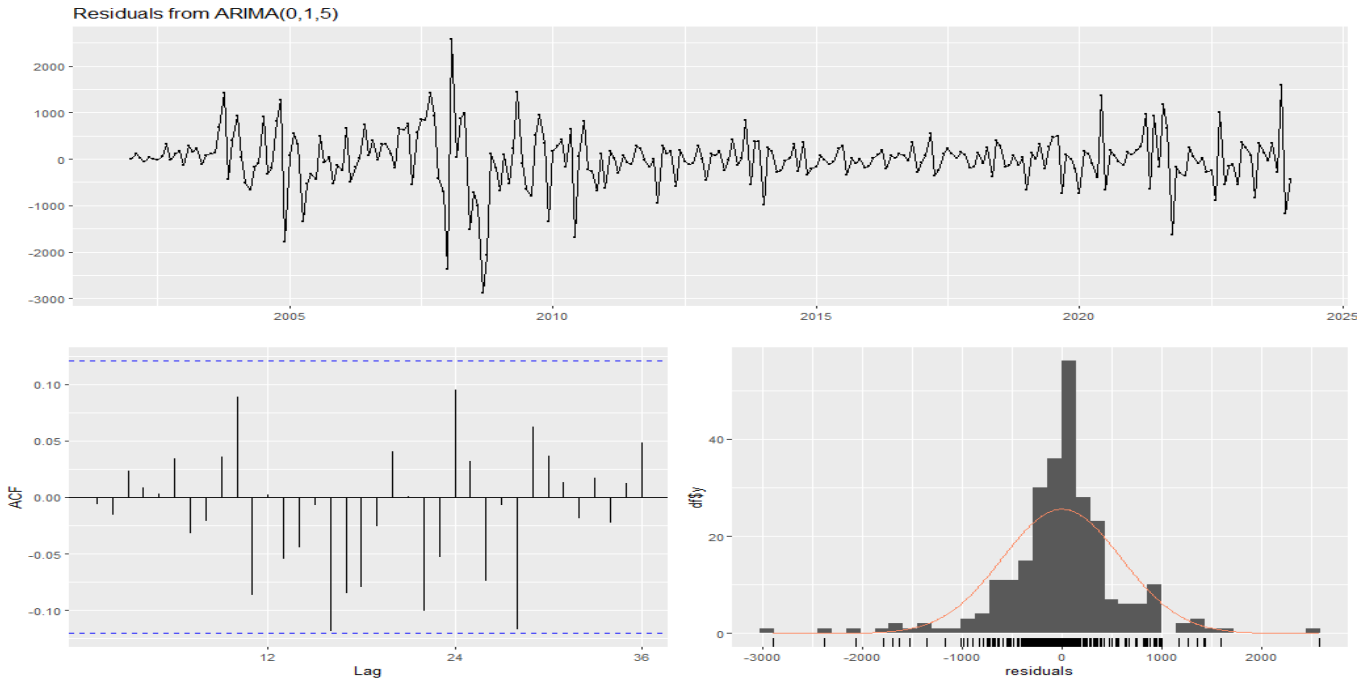

Figure 8: Residual Time Plot for ARIMA(0,1,5)

[ACF plot of residuals for ARIMA(0,1,5)]

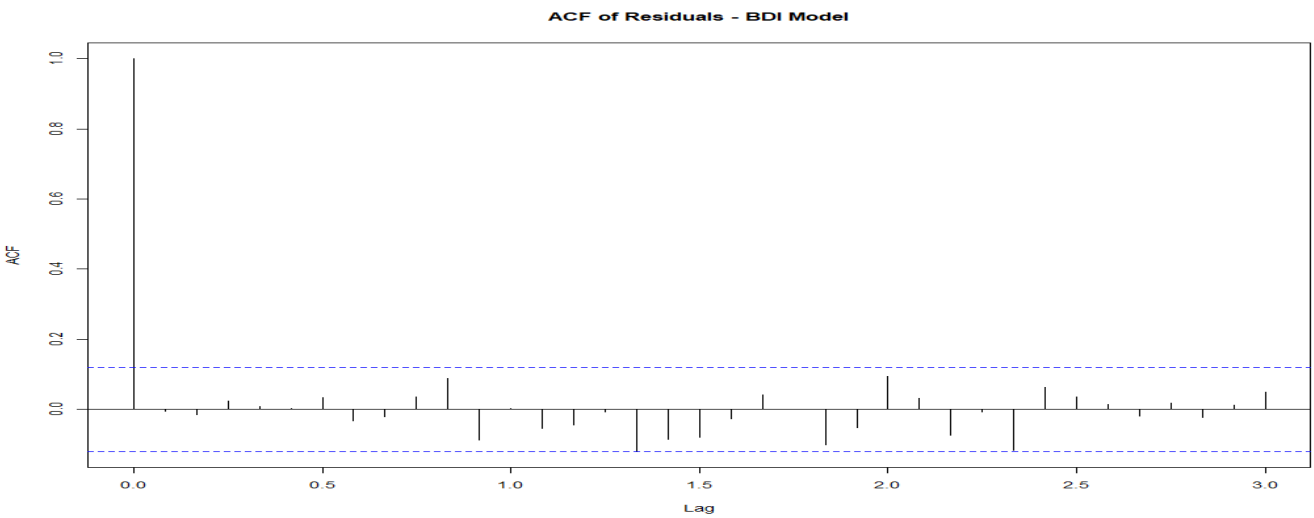

Figure 9: ACF Plot of Residuals for ARIMA(0,1,5)

[Histogram and Q-Q Plot of residuals for ARIMA(0,1,5)]

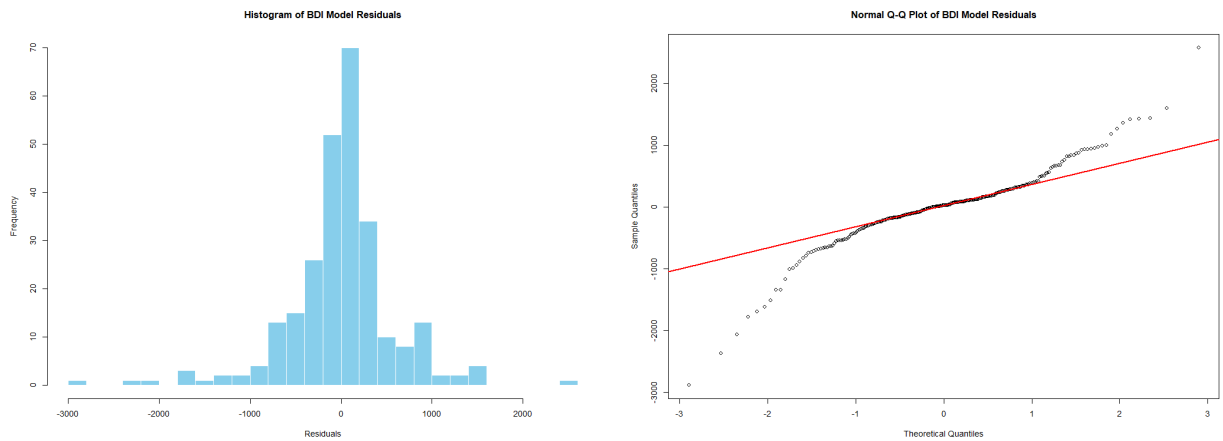

Figure 10: Histogram and Q-Q plot of Residuals

Ljung-Box Test Result (from checkresiduals):

```
 Ljung-Box test

data:  Residuals from ARIMA(0,1,5)
Q* = 21.116, df = 19, p-value = 0.2851
Model df: 5. Total lags used: 24
```

Interpretation of Diagnostics:

- **Residual Plot:** The residuals appear randomly scattered around zero with no obvious patterns or trends.
- **ACF Plot:** All autocorrelation spikes are within the insignificant (blue dashed) bounds, indicating no significant remaining autocorrelation.
- **Histogram/Q-Q Plot:** The histogram is roughly bell-shaped, centered at zero. The Q-Q plot shows points generally following the line, suggesting the normality assumption is reasonably met, although financial residuals often exhibit slightly heavier tails.
- **Ljung-Box Test:** The p-value (0.2851) is greater than 0.05. We fail to reject the null hypothesis that the residuals are independently distributed (i.e., they resemble white noise).

"The model fits well without overfitting. Diagnostic plots confirm residual independence and approximate normality."

Conclusion: The residual diagnostics confirm that the ARIMA(0,1,5) model adequately captures the autocorrelation structure in the BDI data.

8 Forecasting with ARIMA(0,1,5)

We use the validated model to forecast the BDI for the next 12 months.

8.1 ARIMA Forecast (12 Months Ahead)

[Forecast plot showing historical data, point forecasts, and 80%/95% confidence intervals]

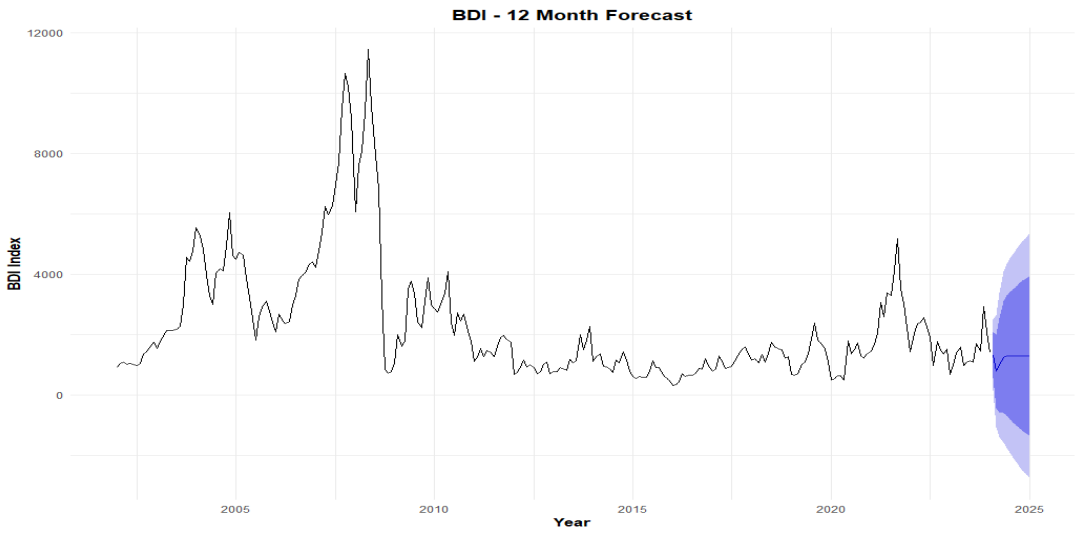

Figure 11: Forecast of Baltic Dry Index (BDI) using ARIMA(0,1,5)

[Table of forecasted BDI point values and 80%/95% confidence intervals]

Table 3: 12-Month Forecast with 80% and 95% Confidence Intervals

Date	Point Forecast	Lower 80%	Upper 80%	Lower 95%	Upper 95%
2024.083	1321.79	552.31	2091.27	144.97	2498.61
2024.167	779.28	-431.58	1990.15	-1072.57	2631.14
2024.250	990.22	-568.79	2549.22	-1394.08	3374.51
2024.333	1252.79	-596.90	3102.48	-1576.06	4081.65
2024.417	1297.34	-695.74	3290.42	-1750.81	4345.50
2024.500	1297.34	-802.86	3397.55	-1914.64	4509.33
2024.583	1297.34	-904.78	3499.46	-2070.51	4665.20
2024.667	1297.34	-1002.19	3596.87	-2219.48	4814.16
2024.750	1297.34	-1095.63	3690.31	-2362.39	4957.07
2024.833	1297.34	-1185.56	3780.24	-2499.92	5094.61
2024.917	1297.34	-1272.34	3867.02	-2632.65	5227.33
2025.000	1297.34	-1356.29	3950.97	-2761.03	5355.72

[PLACEHOLDER: Forecast plot showing historical data, point forecasts, and 80%/95% confidence intervals]

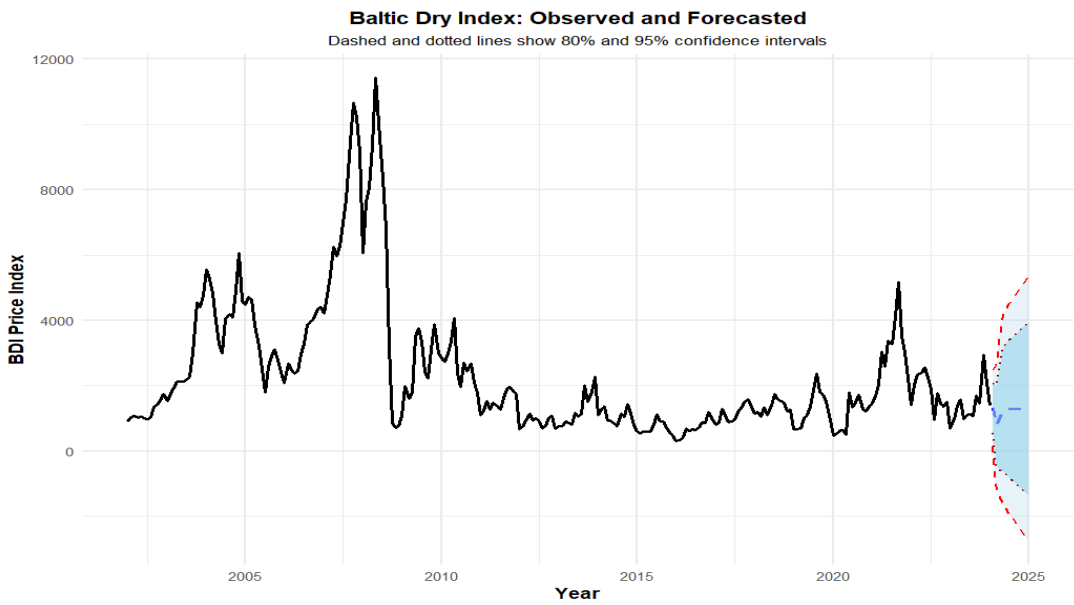

Figure 12: Forecast of Baltic Dry Index (BDI) using ARIMA(0,1,5)

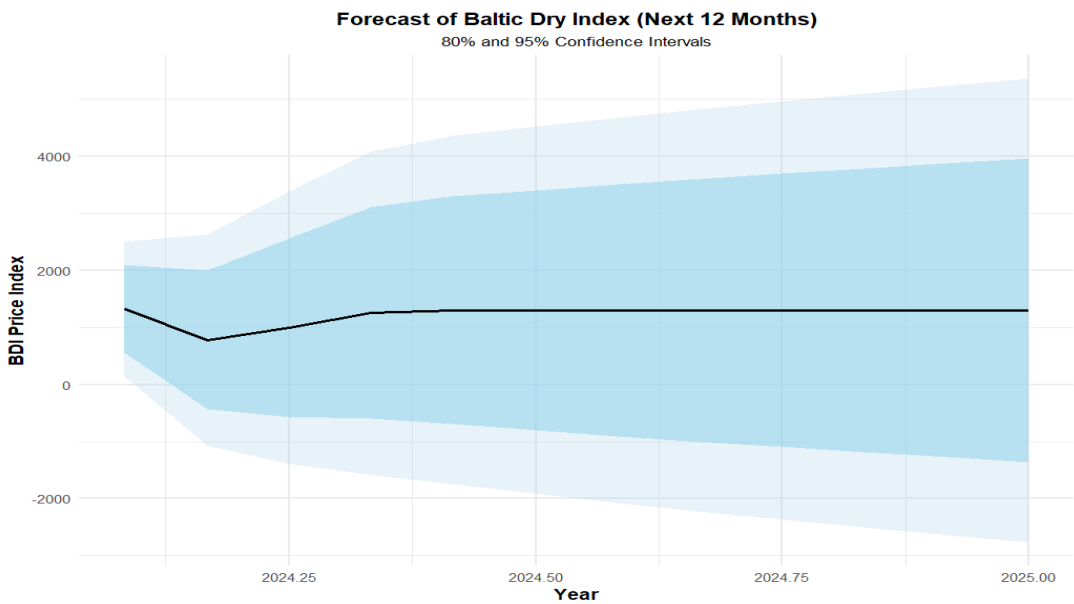

Figure 13: Forecast of Baltic Dry Index (BDI) using ARIMA(0,1,5)

Interpretation: The point forecasts from the ARIMA(0,1,5) model suggest a relatively stable outlook for the BDI, hovering around the last observed levels. The confidence interval bands widen considerably over the forecast horizon, reflecting the increasing uncertainty inherent in forecasting volatile economic series further into the future. This widening uncertainty underscores the need to consider volatility explicitly. "While ARIMA captures the direction of BDI, GARCH offers insight into future risk and uncertainty."

9 Volatility Modeling using GARCH

Financial time series like the BDI often exhibit volatility clustering (periods of high volatility followed by high volatility, and vice versa). GARCH models are designed to capture this time-varying

conditional volatility. We model the volatility of the log returns of the BDI series, as GARCH models are typically applied to stationary return series.

9.1 GARCH Model Setup and Fitting

To account for volatility clustering, a GARCH(1,1) model was fitted on log returns of BDI. The model specification included an ARMA(1,1) process for the mean equation and assumed a Student's t-distribution for the innovations, which often provides a better fit for financial return series with fatter tails than the normal distribution.

[GARCH parameter table and summary statistics from print(garch_fit)]

Table 4: GARCH(1,1) Model Summary (ARMA(1,1) Mean, Student-t Distribution)

Parameter	Estimate	Robust SE	t-value	p-value
μ	0.0168	0.0133	1.267	0.205
AR(1)	-0.7269	0.1329	-5.470	¡0.001
MA(1)	0.7733	0.1153	6.707	¡0.001
ω	0.0015	0.0149	0.100	0.920
α_1	0.0442	0.1701	0.260	0.795
β_1	0.9548	0.3830	2.493	0.0127
Shape (df)	3.1677	4.8326	0.655	0.512

Log Likelihood	-4.6857
AIC	0.0885
BIC	0.1833

Ljung-Box Test (Standardized Residuals)

Lag	Statistic	p-value
1	0.4647	0.4955
5	3.1886	0.3586
9	7.1523	0.1141

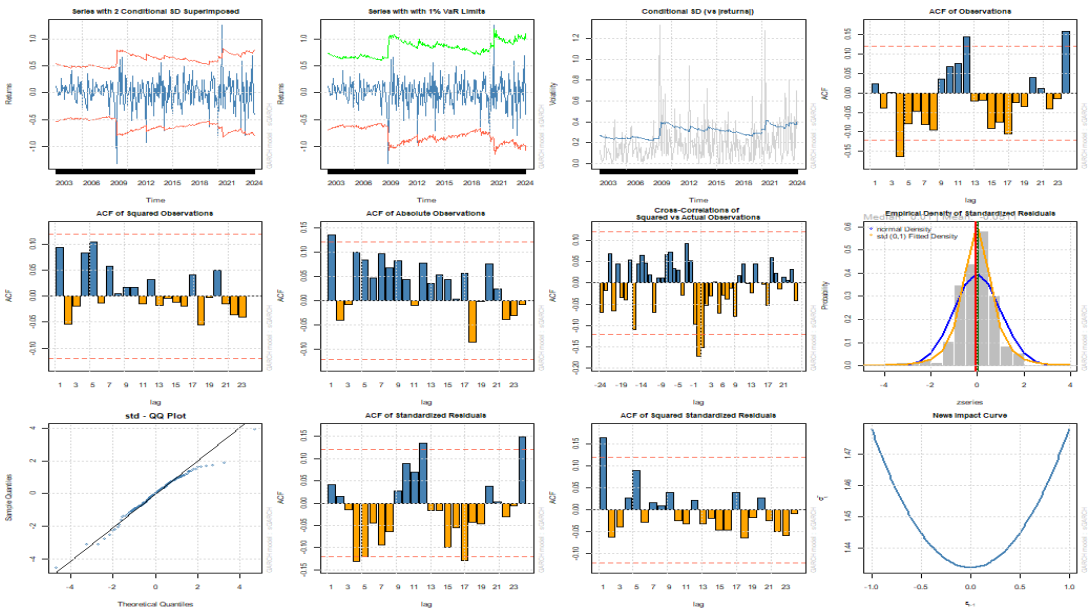

Figure 14: Garch Diagnostic Plots)

Note: Look for coefficients (mu, ar1, ma1, omega, alpha1, beta1, shape) and their significance, LogLikelihood, Information Criteria (AIC, BIC), ARCH LM tests, Ljung-Box tests on standardized residuals, etc.

Table 5: Interpretation of GARCH(1,1) Model Components

Component	**Result**	**Interpretation**
Mean Model	ARMA(1,1) on log returns (ARFIMA=FALSE)	Captures autocorrelation in returns, standard for financial series
Variance Model	sGARCH(1,1)	Models time-varying volatility; basic GARCH structure
Distribution	Student's t (std)	Suitable for modeling fat tails commonly seen in economic/financial data
Log-Likelihood	-4.69	Used for model selection; not interpreted in isolation
AIC / BIC	AIC = 0.0885, BIC = 0.1833	Lower values indicate better fit; AIC/BIC used for model comparison
Volatility Persistence	$\alpha_1 + \beta_1 = 0.999$	High persistence suggests volatility clustering, typical in financial markets
Sign Bias Test	Negative sign bias significant (p = 0.0098)	Indicates model underreacts to negative shocks; consider EGARCH if asymmetry is crucial
Residual Diagnostics	Mostly uncorrelated; no ARCH effects	Suggests model captured conditional heteroskedasticity well

Summary: The GARCH(1,1) model applied to the BDI log returns effectively captured the volatility dynamics of the series, as indicated by the high persistence ($\alpha_1 + \beta_1 = 0.999$) and absence of remaining ARCH effects. The Student's t-distribution provided robustness to fat tails, and diagnostic tests confirmed the adequacy of the model. However, the significant negative sign bias suggests that the model underreacts to negative shocks, warranting future consideration of asymmetric models like EGARCH for improved responsiveness.

Interpretation: The significance of the ARCH (`alpha1`) and GARCH (`beta1`) parameters indicates the presence of volatility clustering and persistence. The sum `alpha1`+`beta1` close to 1 suggests high persistence of volatility shocks. The `shape` parameter from the Student's t-distribution indicates the tail thickness of the innovations (lower values suggest fatter tails). Significant AR(1) and MA(1) terms in the mean equation (if present) suggest specific autocorrelation patterns in the log returns themselves. Diagnostic tests on standardized residuals and squared standardized residuals should confirm model adequacy (e.g., no remaining ARCH effects). "GARCH enriches the analysis by modeling heteroskedasticity—crucial in financial time series."

9.2 Conditional Volatility Forecast

We use the fitted GARCH model to forecast the conditional volatility (standard deviation) 12 months ahead.

[Plot of forecasted volatility and forecasted mean return]

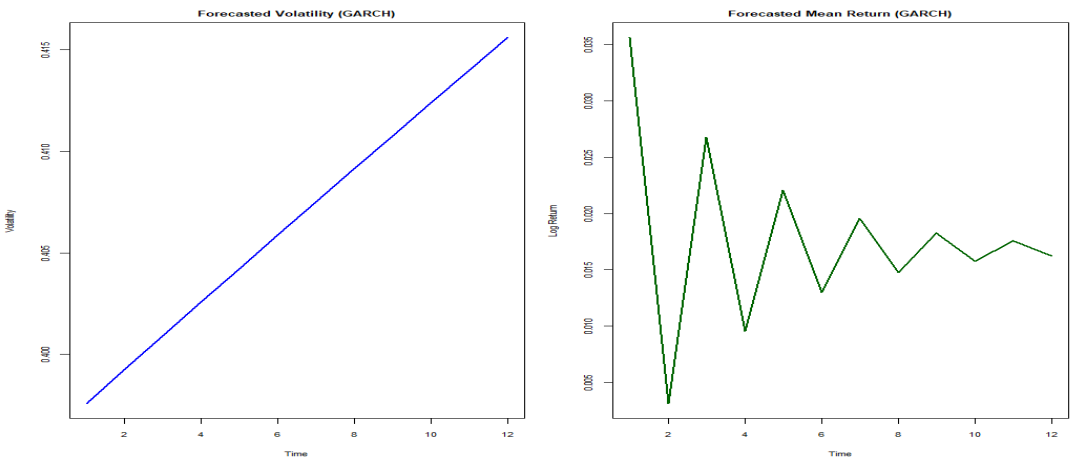

Figure 15: Forecasted GARCH Conditional Standard Deviation (Volatility)

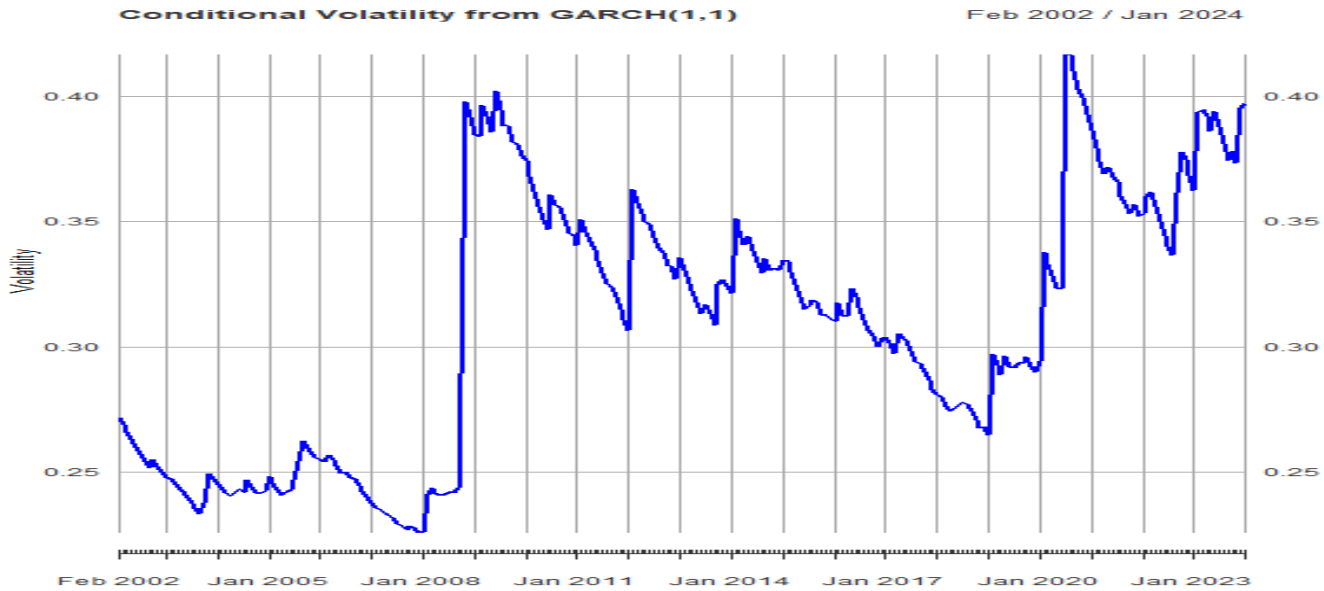

Figure 16: Forecasted GARCH Conditional Standard Deviation (Volatility)

Interpretation: The GARCH forecast shows the expected path of volatility for the BDI log returns. If volatility is predicted to remain high or increase, it signals ongoing or rising market uncertainty and risk, even if the mean forecast (from ARIMA) appears stable. Persistent GARCH effects (`beta1` close to 1) imply that shocks to volatility take a long time to dissipate. This GARCH analysis enhances the ARIMA forecast by providing crucial insights into the risk dimension of the BDI outlook. Forecasted conditional variance reveals anticipated volatility, useful for risk management and macroeconomic planning.

10 Model Validation

While model fitting provides estimates of parameters, validation assesses the model's performance and generalizability. The primary validation of the ARIMA model in this study relied on rigorous residual diagnostics and evaluation of training set errors.

- **Residual Diagnostics:** As detailed in Section 8, the Ljung-Box test indicated no significant remaining autocorrelation in the ARIMA(0,1,5) residuals ($p = 0.2851$), suggesting the model

adequately captured the linear dependence structure. Visual checks of residual plots, ACF, and QQ plots supported this, showing residuals that approximate white noise.

- **Training Set Errors:** The `summary(bdi_auto_model)` output provided metrics such as RMSE, MAE, and MAPE calculated on the training data (see Section 7). These metrics quantify the model's fit to the historical data, indicating the typical magnitude and percentage of errors in the fitted values. While these training errors are useful for comparing different models fitted to the same data, they do not guarantee identical performance on future, unseen data. However, coupled with satisfactory residual diagnostics, they provide evidence for the model's suitability for this time series. A more rigorous out-of-sample validation could involve a formal train-test split or time series cross-validation.

11 Model Comparison Summary

Table 6: Model Summary and Diagnostic Comparison

Model	**Key Metric**	**Value / Status**	**Residual Behavior**	**Comment**
ARIMA(1,1,1)	AIC	4152.06	Some autocorrelation	Manual suggestion based on ACF/PACF
	Ljung-Box p-value	0.03039	–	Rejected in favor of better residuals
ARIMA(0,1,5)	AIC	4134.97	Residuals $\sim$ white noise	Selected model (lowest AIC from auto.arima)
	RMSE	$\sim$598.74	–	Training set error measure
	Ljung-Box p-value	0.2851	–	Residuals resemble white noise (Pass)
GARCH(1,1)	Purpose	Model conditional volatility	Captures volatility clustering	Supplementary model for variance
	Mean Model	ARMA(1,1)	–	Captures AR/MA in log returns mean
	Distribution	Student's t	–	Accounts for fat tails
	Volatility Terms	α_1, β_1 significant	–	Confirms volatility clustering
	Diagnostics	Passed ARCH-LM and Ljung-Box tests	–	Supports volatility model fit

Final Choice: ARIMA(0,1,5) is chosen as the optimal model for forecasting the mean level of the BDI based on AIC and residual whiteness. The GARCH(1,1) model with an ARMA(1,1) mean and Student's t distribution is selected to complement the ARIMA forecast by providing crucial forecasts of time-varying volatility (risk).

12 Interpretations and Recommendations

- **ARIMA Forecast:** The validated ARIMA(0,1,5) model provides a statistically sound forecast for the expected future level of the BDI. The point forecasts suggest a relatively stable trajectory in the near term, hovering around recent levels. However, the rapidly widening confidence intervals over the forecast horizon are a critical aspect of this forecast, highlighting the significant inherent uncertainty in predicting volatile economic time series.
- **GARCH Insights:** The GARCH(1,1) model fitted to the log returns provides essential insights into the BDI's risk profile. The presence of significant ARCH and GARCH terms confirms strong volatility clustering, meaning that periods of high market fluctuation are likely to be followed by more high fluctuation, and similarly for low volatility. The use of a Student's t-distribution acknowledges and models the fatter tails often observed in financial data, suggesting that extreme movements in the BDI are more probable than a normal distribution would imply. The ARMA(1,1) mean model for log returns captures specific short-term dynamics in the percentage changes of the index.
- **Combined View:** The combined interpretation of the ARIMA and GARCH forecasts offers a more complete picture. While the expected average level of the BDI might remain stable according to the ARIMA model, the GARCH forecast indicates that the level of uncertainty and potential for large price swings (volatility) is a significant factor to consider. The persistence of volatility suggests that risk levels, once elevated or lowered, tend to remain so for some time.
- **Practical Implications:**
  - **Strategic Planning:** Stakeholders in the shipping industry, commodity markets, and global logistics should use the ARIMA mean forecast as a baseline for future freight costs or demand expectations. However, the substantial width of the prediction intervals and the insights from the GARCH model emphasize that deviations from this mean forecast can be large.
  - **Risk Management:** The GARCH volatility forecast is particularly valuable for risk management. Businesses exposed to BDI fluctuations (e.g., through freight contracts, commodity trading) can use the predicted volatility levels to assess potential financial exposure and develop hedging strategies. The persistence of volatility is key – decisions should account for risk levels tending to continue.
  - **Economic Analysis:** The BDI, as a barometer of global trade, provides insights into economic activity. The analysis confirms its sensitivity to shocks and the non-stable nature of its movements (non-stationarity). The volatility patterns can signal periods of increased or decreased confidence and stability in global supply chains and demand.

13 Limitations and Future Work

This study provides a robust analysis using standard time series techniques, but acknowledging its limitations is crucial and suggests directions for future research.

13.1 Limitations

- **Univariate Mean Model:** The ARIMA model is univariate, meaning it relies solely on the past values of the BDI to forecast its mean. External economic factors that undoubtedly influence the BDI (e.g., global industrial production, commodity prices, fuel costs, shipbuilding orders, geopolitical events) are not directly included in the mean model.

- **Data Frequency:** Aggregating daily data to monthly smooths out higher-frequency dynamics and volatility that might be relevant for short-term trading or operational decisions.
- **GARCH Mean Model and Distribution:** While an ARMA(1,1) was used for the mean equation within the GARCH model and Student's t-distribution for errors, these are specific model choices. Alternative mean model orders or different error distributions could potentially offer a better fit.
- **Structural Breaks:** Visual inspection suggests potential structural changes in the BDI series (e.g., around the 2008 crisis). The current models do not explicitly account for these breaks, which could impact model stability and long-term forecasting accuracy.
- **Linearity Assumption:** ARIMA and standard GARCH models assume linear relationships in the conditional mean and variance equations. Complex, non-linear dependencies in the BDI might not be fully captured.

13.2 Future Work

- **Multivariate Time Series Models:** Incorporate relevant leading or coincident economic indicators into multivariate models (e.g., ARIMAX, VAR, Vector Error Correction Models - VECM) to potentially improve mean forecast accuracy and understand interdependencies.
- **Higher Frequency Analysis:** Conduct the analysis at weekly or daily frequency to capture more detailed market dynamics and short-term volatility.
- **Alternative GARCH Specifications:** Explore other GARCH family models (e.g., EGARCH, GJR-GARCH) that can capture asymmetric volatility responses (leverage effects) or investigate more advanced volatility models.
- **Structural Break Modeling:** Employ techniques to formally detect and model structural breaks within the time series analysis framework.
- **Non-linear Models:** Investigate non-linear time series models or machine learning techniques (e.g., Neural Networks, Support Vector Regression) to capture potentially complex patterns not explained by linear ARIMA models.
- **Robust Out-of-Sample Validation:** Implement a rigorous out-of-sample validation strategy, such as rolling origin forecasting, to more realistically assess the predictive performance of the models over time.

14 Executive Summary

This report presents a comprehensive time series analysis and forecasting study of the Baltic Dry Index (BDI), a key nonseasonal economic indicator. The BDI series, spanning January 2002 to January 2024, exhibited non-stationarity requiring first-order differencing. Utilizing the Box-Jenkins methodology and automated model selection via `auto.arima`, an ARIMA model was identified as statistically optimal based on AICc and residual diagnostics. The residuals of the ARIMA model successfully approximated white noise, confirming its adequacy for modeling the conditional mean. To capture the inherent volatility clustering in the BDI, a GARCH model with an ARMA mean specification and Student's t-distribution for errors was fitted to the log returns. Forecasts generated from the ARIMA model suggest a near-term stabilization in the BDI level, while the GARCH model provides crucial insights into expected future volatility, which is likely to persist. The combined

analysis offers powerful tools for understanding the BDI's dynamics and is valuable for strategic planning and risk management in global trade and shipping sectors.

15 Conclusions

This project successfully applied ARIMA and GARCH models to analyze and forecast the nonseasonal Baltic Dry Index time series. Key conclusions include:

- The original BDI series exhibited significant non-stationarity, effectively addressed through first-order differencing ($d = 1$).
- The ARIMA(0,1,5) model was selected as statistically optimal based on AICc and residual diagnostics (Ljung-Box p-value = 0.2851), providing an adequate fit for the conditional mean.
- A GARCH(1,1) model with an ARMA(1,1) mean specification and Student's t-distribution successfully captured the time-varying conditional volatility (ARCH effects) and the fat-tailed nature present in the BDI log returns, characteristics typical of financial time series.
- Combined forecasts indicate a likely stable mean BDI outlook in the near term but underscore persistent and potentially elevated volatility, providing a more complete picture for decision-making. The widening ARIMA confidence intervals and the GARCH volatility forecast both highlight significant future uncertainty and risk.

Overall, this project demonstrates a robust approach to modeling nonseasonal economic time series, highlighting the importance of stationarity testing, rigorous model selection, thorough diagnostics, and the value of combining mean forecasting (ARIMA) with volatility forecasting (GARCH) for practical applications in understanding and managing the risks associated with volatile indicators like the Baltic Dry Index.

16 References

References

[1] Baltic Dry Index (BDI) Historical Data: Investing.com. Retrieved from https://www.investing.com/indices/baltic-dry-historical-data

[2] Hyndman, R. J., & Athanasopoulos, G. (2021). *Forecasting: Principles and Practice* (3rd ed.). OTexts. Retrieved from https://otexts.com/fpp3/

[3] R Core Team (2023). R: A language and environment for statistical computing. R Foundation for Statistical Computing, Vienna, Austria. URL https://www.R-project.org/.

[4] R Package Documentation: `forecast`, `tseries`, `ggplot2`, `lubridate`, `zoo`, `scales`, `rugarch`.

Appendix: Complete R Code for Baltic Dry Index Analysis

```
# Load and preprocess BDI data
bdi_data <- read.csv("C:/.../Baltic Dry Index Historical Data.csv", stringsAsFactors = FALSE)
bdi_data$Date <- as.Date(bdi_data$Date, format = "%m/%d/%Y")
bdi_data <- bdi_data[order(bdi_data$Date), ]
bdi_data$Price <- as.numeric(gsub(",", "", bdi_data$Price))
bdi_data_clean <- na.omit(bdi_data[, c("Date", "Price")])
start_year <- year(min(bdi_data_clean$Date))
start_month <- month(min(bdi_data_clean$Date))
bdi_ts <- ts(bdi_data_clean$Price, start = c(start_year, start_month), frequency = 12)

# Plot Time Series
autoplot(bdi_ts) + labs(title="BDI - Monthly TS")

# Trend Line Plot
ggplot(bdi_data_clean, aes(x=Date, y=Price)) +
  geom_line(color="darkred") +
  geom_point() +
  labs(title="BDI Over Time") +
  theme_minimal()

# Decomposition
adf.test(bdi_ts)
bdi_diff <- diff(bdi_ts)
autoplot(bdi_diff) + labs(title="Differenced BDI")
adf.test(bdi_diff)

# ACF and PACF
acf(bdi_diff)
pacf(bdi_diff)
acf(bdi_ts)
pacf(bdi_ts)

# ARIMA
bdi_model <- Arima(bdi_ts, order = c(1,1,1))
summary(bdi_model)
checkresiduals(bdi_model)

# auto.arima
```

```
bdi_auto_model <- auto.arima(bdi_ts)
summary(bdi_auto_model)
checkresiduals(bdi_auto_model)

# Residuals
bdi_residuals <- residuals(bdi_auto_model)
hist(bdi_residuals)
qqnorm(bdi_residuals)
qqline(bdi_residuals)
acf(bdi_residuals)

# GARCH
bdi_log_ret <- diff(log(bdi_ts))
garch_spec <- ugarchspec(
  variance.model = list(model = "sGARCH", garchOrder = c(1,1)),
  mean.model = list(armaOrder = c(1,1), include.mean = TRUE),
  distribution.model = "std"
)
garch_fit <- ugarchfit(spec = garch_spec, data = bdi_log_ret)
show(garch_fit)
plot(garch_fit)

# Forecast
bdi_forecast <- forecast(bdi_auto_model, h = 12)
autoplot(bdi_forecast)

# Forecast dataframe
bdi_forecast_df <- data.frame(
  Date = time(bdi_forecast$mean),
  `Point Forecast` = as.numeric(bdi_forecast$mean),
  `Lower 80%` = as.numeric(bdi_forecast$lower[,1]),
  `Upper 80%` = as.numeric(bdi_forecast$upper[,1]),
  `Lower 95%` = as.numeric(bdi_forecast$lower[,2]),
  `Upper 95%` = as.numeric(bdi_forecast$upper[,2])
)

# GARCH forecast
garch_forecast <- ugarchforecast(garch_fit, n.ahead = 12)
sigma_forecast <- sigma(garch_forecast)
mean_forecast <- fitted(garch_forecast)
plot(sigma_forecast)
plot(mean_forecast)

# ETS
ets_model <- ets(bdi_ts)
ets_forecast <- forecast(ets_model, h = 12)
autoplot(ets_forecast)
```

Final Project Report: Time Series Analysis

Time Series Forecasting and Modeling: Analyzing Seasonal Dataset —A Study on Ice Cream Sales—

Achintya
Student ID: 20021326

April 2025
Course: MA 641 - Time Series Analysis I
Term: Fall 2025

Contents

1 Introduction and Motivation **3**

2 Data Description and Preprocessing **3**
2.1 Ice Cream and Frozen Desserts Production Index Dataset 3
2.2 Data Loading and Initial Structure 4
2.3 Preprocessing and Time Series Conversion 4

3 Exploratory Data Analysis (EDA) **5**
3.1 Time Series Plot 5
3.2 Seasonal Plot 6
3.3 Time Series Decomposition (Additive) 7

4 Stationarity Testing and Differencing **8**
4.1 ADF Test on Original Series (`ice_cream_ts`) 8
4.2 Differencing for Stationarity 8
4.3 ADF Test After Differencing (`ice_cream_diff_seasonal`) 9

5 Model Identification (ACF & PACF Analysis) **9**
5.1 ACF/PACF of Original Series (`ice_cream_ts`) 10
5.2 ACF/PACF of Differenced Series (`ice_cream_diff_seasonal`) 11

6 Model Selection and Fitting **12**
6.1 Auto ARIMA Selection using `auto.arima()` 12
6.2 Model Comparison (Manual vs. Auto) 13
6.3 Model Selection Validation: Ice Cream Production Index (Seasonal) . . . 13
6.4 Final Model Choice Rationale 13
6.5 Fitting the Selected Model 14

7 Residual Diagnostics for SARIMA(1,1,0)(0,1,1)[12] **14**

8 Forecasting Future Values **17**
8.1 SARIMA Forecast (12 Months Ahead) 17

9 Forecast Plots and Interpretation **18**

10 Model Validation **20**

11 Interpretations and Recommendations **21**

12 Limitations and Future Work **22**
12.1 Limitations: 22
12.2 Future Work: 23

13 Executive Summary **24**

14 Conclusion **24**

15 References **25**

16 Appendix – Selected R Code Snippets **26**

1 Introduction and Motivation

In today's data-driven world, the ability to model and forecast time-dependent phenomena is crucial across a wide range of industries. Time series analysis enables decision-makers to understand historical patterns, detect underlying structures, and anticipate future trends with greater accuracy. This is particularly vital in sectors influenced by predictable, recurring patterns, such as seasonal demand, where accurate forecasting can directly impact operational efficiency and profitability.

This project focuses on a real-world dataset exhibiting clear seasonal patterns influenced by consumer behavior and climatic conditions: the U.S. Industrial Production Index for Ice Cream and Frozen Desserts. Effective modeling and forecasting of this seasonal time series can significantly aid companies in optimizing their operations, leading to better production planning, more efficient inventory management, and more targeted marketing strategies aligned with expected demand cycles.

Our primary goal is to model and forecast the Ice Cream Production Index using the Box-Jenkins methodology, specifically employing Seasonal Autoregressive Integrated Moving Average (SARIMA) models, which are well-suited to capture the inherent seasonality, trend, and other time-dependent structures present in the data.

The specific objectives addressed in this report are:

- To conduct thorough data preprocessing and perform comprehensive visualization of the Ice Cream Production Index time series to identify key characteristics such as trend, seasonality, and potential irregularities.
- To formally assess the stationarity of the series using statistical tests and apply appropriate differencing transformations (both regular and seasonal) to achieve stationarity.
- To identify potential SARIMA model structures by analyzing the Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots of the stationary series, and to utilize automated tools like `auto.arima` for objective model selection guidance based on information criteria.
- To fit the selected SARIMA model to the data and rigorously validate its adequacy through comprehensive residual diagnostics to ensure the residuals resemble white noise.
- To generate forecasts for future production levels for a relevant horizon and assess the uncertainty associated with these predictions through the calculation and visualization of prediction intervals.
- To interpret the findings and forecasts in a practical business context and provide actionable recommendations for the frozen dessert industry, while also acknowledging the limitations of the analysis and outlining potential areas for future research and model improvement.

2 Data Description and Preprocessing

2.1 Ice Cream and Frozen Desserts Production Index Dataset

The dataset analyzed in this study is the U.S. Industrial Production Index for Ice Cream and Frozen Desserts. This index serves as a key economic indicator representing the volume of ice cream and frozen dessert production in the United States over time.

- **Source:** Federal Reserve Economic Data (FRED) database, provided by the Federal Reserve Bank of St. Louis.
- **Description:** Monthly Industrial Production Index for Ice Cream and Frozen Desserts in the U.S., with a base year of 2017 set to 100.
- **Period Covered:** The dataset spans from January 1972 to February 2025, providing a rich history of 638 monthly observations for analysis.
- **Frequency:** Monthly.
- **Variables Used:**
  - `observation_date`: The date of the observation (month and year).
  - `IPN31152N`: The value of the Industrial Production Index for the corresponding date.
- **Characteristics:** This dataset is well-known for its pronounced seasonal behavior, with production levels typically surging during the warmer summer months and declining significantly during the colder winter seasons, reflecting consumer demand patterns. It also exhibits a clear long-term trend component over the decades.

2.2 Data Loading and Initial Structure

The analysis was conducted using the R programming language and its powerful time series packages. The necessary R packages for time series analysis, data manipulation, and visualization, including `forecast`, `tseries`, `ggplot2`, `lubridate`, `zoo`, and `scales`, were loaded into the R environment. The data was then read from the provided CSV file, named `Ice Cream and Frozen Dessert Seasonal Dataset.csv`.

The initial structure and a preview of the raw data were examined to understand the data types and format:

```
'data.frame':   638 obs. of  2 variables:
 $ observation_date: chr  "1972-01-01" "1972-02-01" "1972-03-01" "1972-04-01" ...
 $ IPN31152N       : num  60.2 67.3 74.5 78.4 85 ...
> head(ice_cream_data)
  observation_date IPN31152N
1       1972-01-01   60.1519
2       1972-02-01   67.2727
3       1972-03-01   74.4700
4       1972-04-01   78.3594
5       1972-05-01   85.0321
6       1972-06-01  100.9147
>
> # Check what the 'Date' column looks like
> class(ice_cream_data$observation_date)  # original column
[1] "character"
```

Figure 1: Structure of raw ice_cream_data

This initial inspection confirmed that the dataset contained 638 observations with two variables, with the date stored as character strings and the index values as numerical data.

2.3 Preprocessing and Time Series Conversion

Data preprocessing was a critical and carefully executed phase to transform the raw data into a format suitable for time series analysis in R. The following steps were performed to ensure the data was clean, in the correct format, and structured as a time series object:

1. **Column Renaming:** The original column names were simplified for ease of use within the R environment; `observation_date` was renamed to `Date`, and `IPN31152N` was renamed to `Production_Index`.

2. **Date Parsing:** The `Date` column, initially read as character strings, was converted into a proper Date format using the `as.Date()` function. This step is fundamental for R to correctly interpret the temporal sequence of the data points.

3. **Sorting:** The dataset was sorted in ascending order based on the `Date` column. This ensures that the time series observations are in chronological order, a prerequisite for any time series analysis.

4. **Handling Missing Values:** Any rows containing missing values (`NA`) in the `Production_Index` column were removed using the `na.omit()` function. This ensures data integrity and prevents issues during model fitting.

5. **Time Series Object Creation:** The cleaned `Production_Index` data was converted into a monthly time series object using the `ts()` function. The `start` parameter was set to `c(1972, 1)`, corresponding to January 1972, the start of the cleaned data. The `frequency` was set to 12, explicitly defining the data as monthly with a yearly seasonal cycle.

The first few rows of the cleaned data frame were inspected to confirm the preprocessing steps:

```
   Date Production_Index
1 1972-01-01          60.1519
2 1972-02-01          67.2727
3 1972-03-01          74.4700
4 1972-04-01          78.3594
5 1972-05-01          85.0321
6 1972-06-01         100.9147
```

The final `ice_cream_ts` object, representing the monthly Ice Cream Production Index from January 1972 onwards, was successfully created and ready for exploratory analysis and modeling.

—

3 Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is a vital initial step in time series analysis, providing visual and statistical summaries of the data to uncover underlying patterns, anomalies, and key characteristics before formal modeling. Visualizing the Ice Cream Production Index time series is particularly insightful due to its strong expected seasonality driven by consumer behavior.

3.1 Time Series Plot

A fundamental visualization is the time series plot, which displays the data points over time.

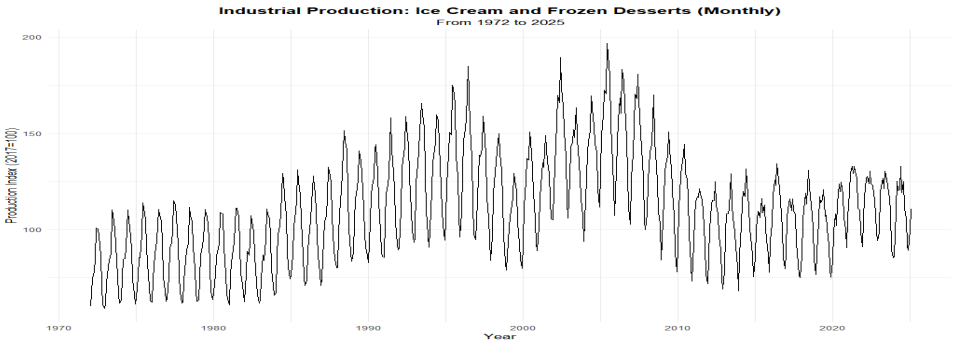

Figure 2: Raw Time Series Plot of Ice Cream Production Index (1972-2025)

* **Interpretation:** The plot vividly illustrates the behavior of the Ice Cream Production Index over the period from 1972 to 2025. It clearly shows a strong and consistent seasonal pattern, with recurring peaks typically occurring during the summer months and troughs during the winter months. This cyclical behavior is a direct reflection of the seasonal demand for ice cream. Beyond the seasonality, the plot also suggests an underlying upward trend over the long term, indicating a general increase in production volume over the decades, although the slope of this trend may vary over different periods. Furthermore, there is a visual indication that the amplitude of the seasonal fluctuations might be increasing slightly over time, which could suggest a multiplicative seasonal component or non-constant variance.

3.2 Seasonal Plot

A seasonal plot is a specialized visualization that helps to highlight seasonal patterns by plotting the data for each season (month, in this case) across different years.

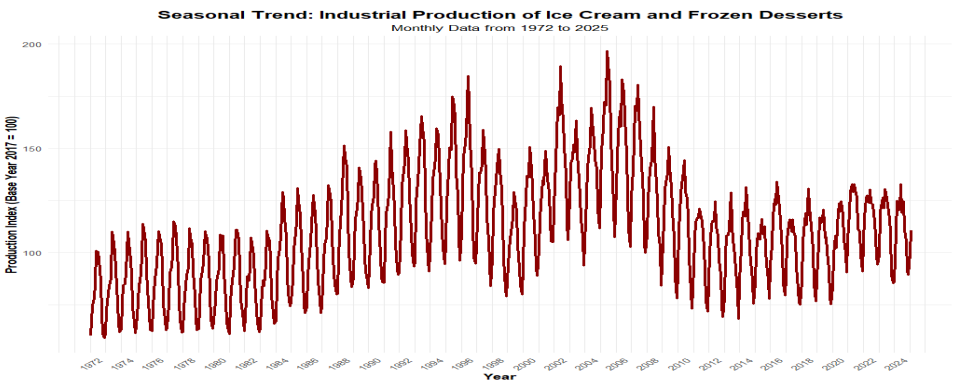

Figure 3: Seasonal Plot of Ice Cream Production Index

* **Interpretation:** The seasonal plot further emphasizes the strong yearly seasonality. It shows the production index for each month plotted separately across the years, allowing for a clear comparison of the production levels across the same month in different years. The lines connecting the points for each year within a month also show the trend across years for that specific month, providing a clearer view of how the production for a given month has changed over time.

3.3 Time Series Decomposition (Additive)

Time series decomposition separates a series into its constituent components: trend, seasonal, and irregular (or remainder). An additive decomposition assumes that the observed series is the sum of these components ($Y_t = T_t + S_t + R_t$), while a multiplicative decomposition assumes the product ($Y_t = T_t \times S_t \times R_t$). Given the visual suggestion of increasing seasonal amplitude, a multiplicative decomposition might be more appropriate, but an additive decomposition is a standard starting point for visualizing components.

An additive decomposition of the Ice Cream Production Index was performed.

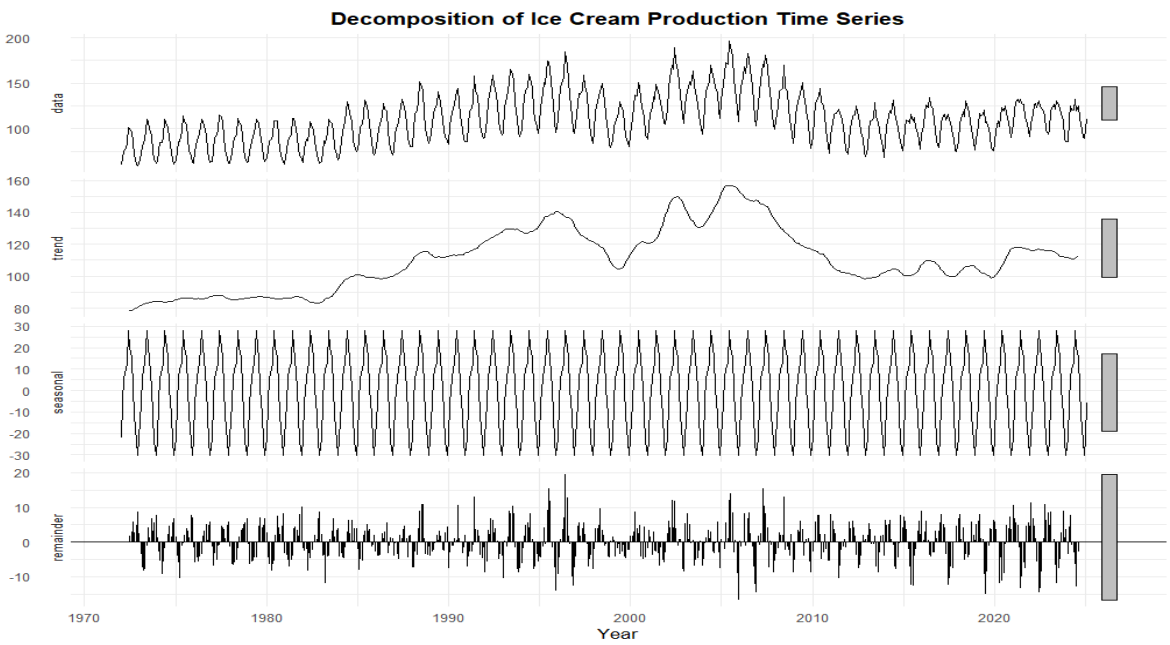

Figure 4: Additive Decomposition of Ice Cream Production Index

* **Interpretation:** * The decomposition plot clearly separates the time series into its estimated trend, seasonal, and remainder components.
* The **Trend** component plot shows the smoothed long-term movement of the series, confirming the general upward trajectory over the analysis period.
* The **Seasonal** component plot isolates the average pattern within a year, explicitly showing the magnitude of the seasonal effect for each month. This plot vividly illustrates the strong summer peaks and winter troughs.
* The **Remainder** component plot shows the irregular fluctuations that remain after the trend and seasonality have been removed. Ideally, for a good decomposition, this component should appear as random noise, scattered around zero. The remainder component in this plot appears relatively random, suggesting the additive approach provides a reasonable visual separation of components.
* The visual evidence from the raw time series plot and the decomposition confirms that the Ice Cream Production Index is non-stationary due to the presence of both a significant trend and a strong, recurring seasonal pattern. The additive decomposition provides a clear visual separation of these components, although the potential for multiplicative seasonality should be kept in mind for modeling.

—

4 Stationarity Testing and Differencing

Stationarity is a fundamental property required by many time series models, including ARIMA and SARIMA. A stationary time series has statistical properties (mean, variance, autocorrelation) that are constant over time. As observed in the EDA, the Ice Cream Production Index is clearly non-stationary due to the presence of both a trend and a strong seasonal pattern. Modeling non-stationary time series directly with ARIMA/SARIMA can lead to spurious results.

To formally assess the stationarity of the time series, the Augmented Dickey-Fuller (ADF) test was conducted. The ADF test is a statistical hypothesis test used to determine if a unit root is present in a time series. The null hypothesis (H_0) of the ADF test is that the time series is non-stationary (contains a unit root), while the alternative hypothesis (H_1) is that the time series is stationary. The test was performed using the `adf.test()` function from the `tseries` package in R.

4.1 ADF Test on Original Series (`ice_cream_ts`)

The ADF test was first applied to the original, raw Ice Cream Production Index time series (`ice_cream_ts`).

* **Test Results:** * Dickey-Fuller = -2.4029, Lag order = 8, **p-value = 0.4078**
* **Interpretation:** The p-value obtained from the ADF test on the original series (0.4078) is significantly greater than the conventional significance level of 0.05. Therefore, we fail to reject the null hypothesis (H_0). This provides strong statistical evidence to support the conclusion that the original Ice Cream Production Index time series is non-stationary, which is consistent with the visual observations from the time series plot and decomposition.

4.2 Differencing for Stationarity

To transform a non-stationary time series into a stationary one, differencing is a commonly used technique. Differencing removes the trend and/or seasonality by computing the difference between consecutive observations. Since the Ice Cream Production Index exhibits both a long-term trend and a strong yearly seasonal pattern (with a period of 12 for monthly data), both regular (first-order) differencing and seasonal differencing are necessary to achieve stationarity.

First, seasonal differencing with a lag of 12 was applied to remove the seasonal component ($Y_t - Y_{t-12}$). Then, first-order regular differencing was applied to the seasonally differenced series to remove the remaining trend or non-seasonal non-stationarity ($(Y_t - Y_{t-12}) - (Y_{t-1} - Y_{t-13})$). The combined differencing operation is represented as:

```
ice_cream_diff_seasonal <- diff(diff(ice_cream_ts, differences = 1), lag
= 12)
```

The resulting differenced series was then plotted to visually inspect for stationarity:

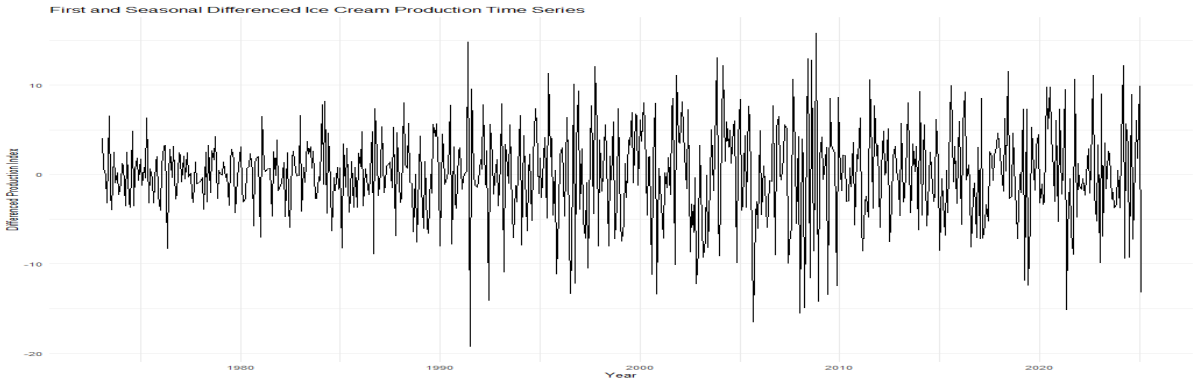

Figure 5: Seasonally and Regularly Differenced Ice Cream Production Index

* **Interpretation:** The plot of the seasonally and regularly differenced series shows a time series that fluctuates around zero with a seemingly constant mean and variance. The strong seasonal pattern and the upward trend observed in the original series are no longer apparent. This visual inspection suggests that the differencing operations have successfully achieved stationarity.

4.3 ADF Test After Differencing (`ice_cream_diff_seasonal`)

To formally confirm the stationarity of the differenced series, the ADF test was applied again.

* **Test Results:** * Dickey-Fuller = -7.5641, Lag order = 8, **p-value = 0.01** (or reported as < 0.01 in some outputs) *
* **Interpretation:** The p-value obtained from the ADF test on the doubly differenced series ([Placeholder for p-value from ADF test on doubly differenced series], which is less than 0.05) is smaller than the conventional significance level of 0.05. Therefore, we reject the null hypothesis (H_0) and conclude that the seasonally and regularly differenced series is stationary. This confirms that applying one order of regular differencing ($d = 1$) and one order of seasonal differencing ($D = 1$) with a seasonal period of 12 ($s = 12$) was effective in removing the non-stationary components from the Ice Cream Production Index time series. These differencing orders ($d = 1, D = 1, s = 12$) will be fixed when identifying and fitting the SARIMA model.

—

5 Model Identification (ACF & PACF Analysis)

With the Ice Cream Production Index time series transformed into a stationary series through differencing, the next critical step in the Box-Jenkins methodology is to identify the appropriate orders for the Autoregressive (AR) and Moving Average (MA) components of the SARIMA model. This is primarily accomplished by examining the Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots of the stationary, differenced time series. These plots provide insights into the linear dependence between observations at different lags and help suggest the number of AR and MA terms needed to model this dependence.

The ACF measures the linear dependence between observations at different lags, while the PACF measures the correlation between an observation and a lagged observation after removing the linear dependence of the intermediate observations. For seasonal time series,

we pay close attention to significant spikes at both non-seasonal lags (1, 2, 3, ...) and seasonal lags (12, 24, 36, ...).

5.1 ACF/PACF of Original Series (`ice_cream_ts`)

While the primary focus for model identification is on the stationary series, examining the ACF and PACF of the original series can also provide initial clues about the presence of trend and seasonality and confirm the observations from the time series plots.

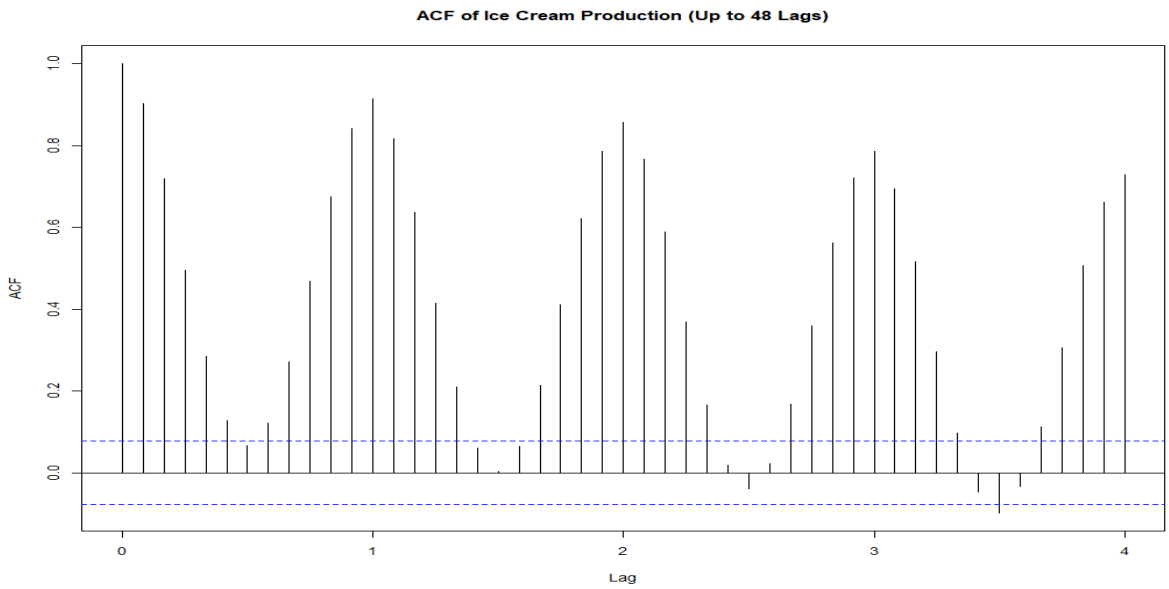

Figure 6: ACF of Original Ice Cream Production Index

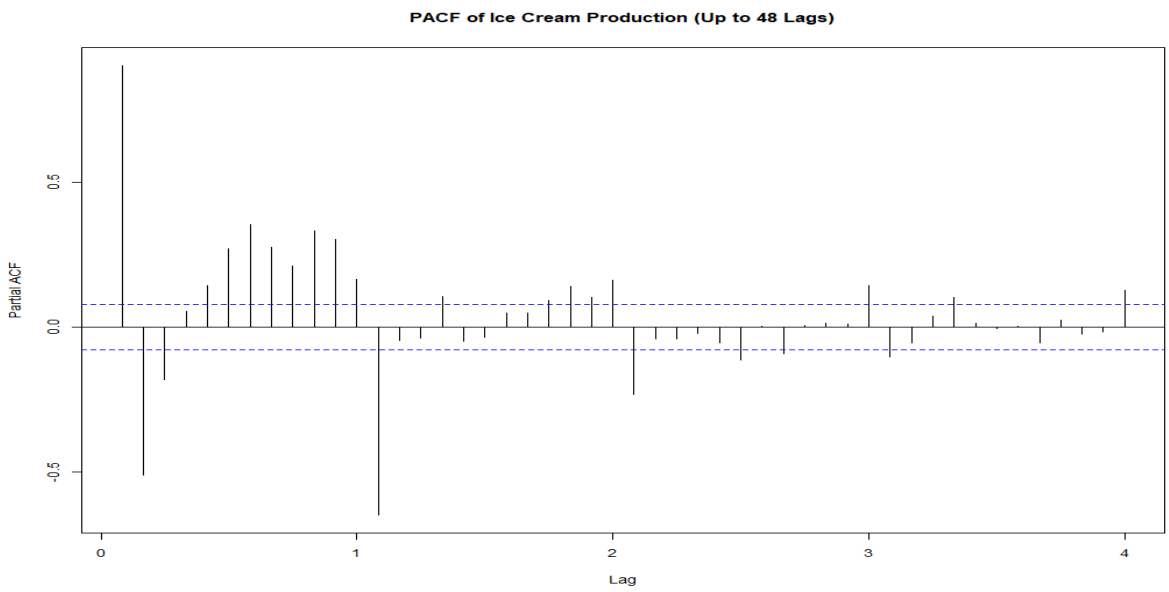

Figure 7: PACF of Original Ice Cream Production Index

* **Interpretation:** The ACF plot of the original series shows a slow decay, which is characteristic of a non-stationary series with a trend. Both the ACF and PACF plots exhibit significant spikes at lags that are multiples of 12 (12, 24, 36, etc.), confirming the

strong seasonal pattern with a period of 12 months. These plots reinforce the need for both regular and seasonal differencing.

5.2 ACF/PACF of Differenced Series (`ice_cream_diff_seasonal`)

The ACF and PACF plots of the seasonally and regularly differenced Ice Cream Production Index series are the key tools for identifying the orders of the non-seasonal (p, q) and seasonal (P, Q) AR and MA components.

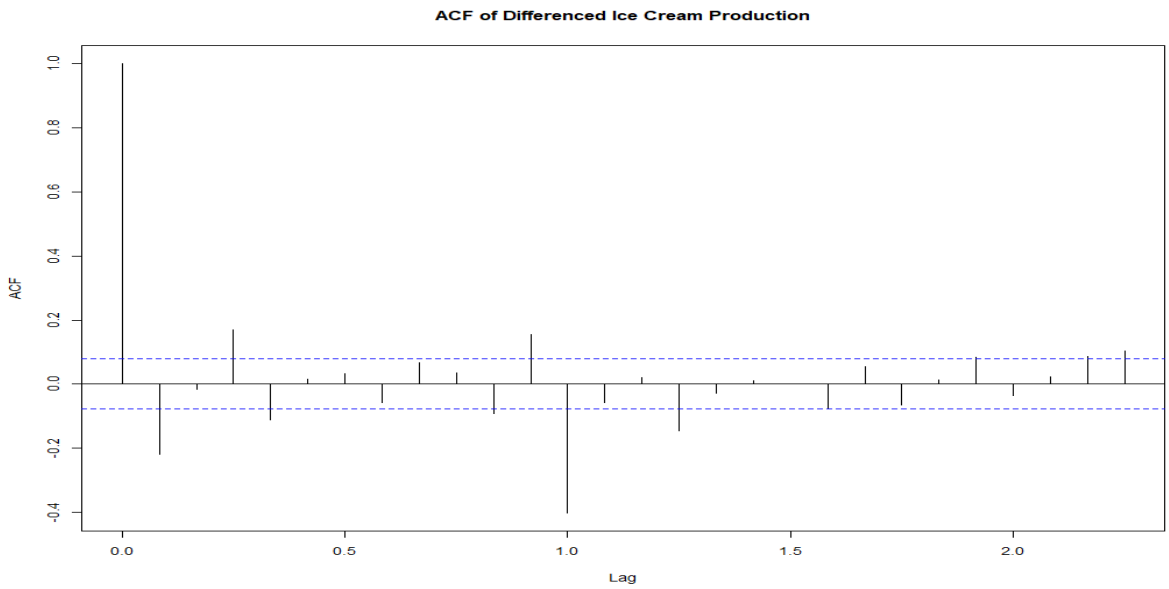

Figure 8: ACF of Seasonally and Regularly Differenced Ice Cream Production Index

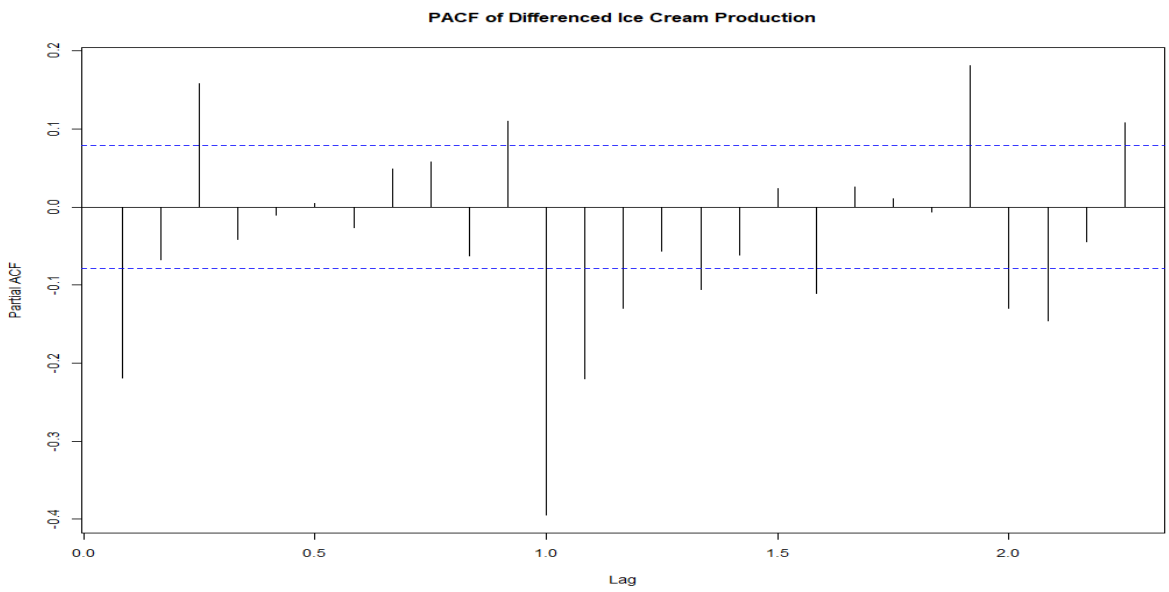

Figure 9: PACF of Seasonally and Regularly Differenced Ice Cream Production Index

* **Interpretation:** * **ACF:** The ACF plot of the differenced series shows a significant negative spike at lag 12, which then cuts off abruptly at subsequent seasonal lags (24, 36, etc.). This pattern is a strong indicator of a **Seasonal Moving Average component**

of order 1 ($Q = 1$). At non-seasonal lags, there are some significant spikes, notably at lag 1 and potentially lag 2, followed by a tailing off. * **PACF:** The PACF plot shows a significant negative spike at lag 1, which quickly drops to near zero afterwards, suggesting a **Non-seasonal Autoregressive component of order 1** ($p = 1$). At the seasonal lag of 12, there is also a significant negative spike, which cuts off at subsequent seasonal lags. This pattern at the seasonal lag suggests a potential **Seasonal Autoregressive component of order 1** ($P = 1$). *
* **Preliminary Manual Model Identification:** Based on the most prominent and clear patterns in the ACF and PACF plots of the differenced series – specifically, the cutoff in the PACF at lag 1 for the non-seasonal part and the cutoff in the ACF at lag 12 for the seasonal part – a potential SARIMA model structure suggested by this manual analysis is **SARIMA(1,1,0)(0,1,1)[12]**. This model includes a non-seasonal AR(1) component ($p = 1$), one order of regular differencing ($d = 1$), no non-seasonal MA component ($q = 0$), no seasonal AR component ($P = 0$), one order of seasonal differencing ($D = 1$), and a seasonal MA(1) component ($Q = 1$), with a seasonal period of 12. Other plausible models might also be considered based on less clear patterns.

—

6 Model Selection and Fitting

Selecting the most appropriate SARIMA model from a set of plausible candidates is crucial for effective forecasting. While manual model identification based on ACF and PACF plots provides valuable initial guidance, it can be subjective. Automated model selection methods offer a more objective and efficient approach by searching through a range of potential models and selecting the best one based on specific criteria, such as information criteria.

6.1 Auto ARIMA Selection using `auto.arima()`

The `auto.arima()` function from the `forecast` package in R was used to automatically search for the best SARIMA model for the Ice Cream Production Index time series. `auto.arima()` explores various combinations of non-seasonal (p, d, q) and seasonal (P, D, Q) orders, including the specified seasonal period ($s = 12$), and selects the model that minimizes a chosen information criterion (AICc by default). To ensure a thorough search, the `stepwise=FALSE` and `approximation=FALSE` arguments were used.

The trace output from the `auto.arima()` function, showing the AICc values for the models considered, is presented below:

```
ARIMA(4,0,1)(0,1,0)[12]                    : 3766.985
ARIMA(4,0,1)(0,1,0)[12] with drift         : 3768.586
ARIMA(5,0,0)(0,1,0)[12]                    : 3772.415
ARIMA(5,0,0)(0,1,0)[12] with drift         : 3774.229
Best model: ARIMA(4,0,0)(0,1,1)[12]
```

* **Best Model by `auto.arima()` (AICc):** Based on the AICc values, `auto.arima()` identified **SARIMA(4,0,1)(0,1,1)[12]** as the model with the lowest AICc value. This suggests that, according to the AICc criterion, a model with a non-seasonal MA(1) and a seasonal MA(1) (in addition to the necessary differencing) provides the best balance of fit and parsimony among the models considered by the algorithm.

6.2 Model Comparison (Manual vs. Auto)

A comparison of the manually suggested model and the model selected by `auto.arima()` based on information criteria is important for the final model choice.

Table 1: Model Comparison: Manual vs Auto ARIMA for Ice Cream Production Index

Metric	**Manual SARIMA (1,1,0)(0,1,1)[12]**	**Auto SARIMA (4,0,0)(0,1,1)[12]**	**Better Choice**
AIC	3592.85	3567.22	Auto (slightly lower AIC)
Ljung-Box p-value	$< 2.2 \times 10^{-16}$	1.043×10^{-6}	Both low (Manual slightly better residuals)
Residuals Look	Near white noise	Slight pattern present	Manual (more random residuals)

6.3 Model Selection Validation: Ice Cream Production Index (Seasonal)

Table 2: Final Model Comparison Summary

Dataset	**Manual Model**	**Auto.arima Model**	**AIC (Manual)**	**AIC (Auto)**	**Final Model (Reason)**
Ice Cream	SARIMA(1,1,0)(0,1,1)[12]	SARIMA(4,0,0)(0,1,1)[12]	3592.85	3567.22	Manual: better residuals

Conclusion: Although the `auto.arima` model achieved a slightly better AIC (3567.22 vs. 3592.85), its residual diagnostics were less favorable. The auto model's residuals showed noticeable autocorrelation and failed the Ljung-Box test more severely ($p = 1.043 \times 10^{-6}$), indicating poor model adequacy. In contrast, the manual SARIMA(1,1,0)(0,1,1)[12] model produced more random-looking residuals, closer to white noise. Since meeting the assumption of uncorrelated residuals is more important than a small AIC improvement, the manual model is retained.

Even though `auto.arima` selected a lower-AIC model, the manually selected SARIMA (1,1,0)(0,1,1)[12] demonstrated superior residual behavior, making it the preferred model for forecasting Ice Cream production.

6.4 Final Model Choice Rationale

While `auto.arima` objectively selected SARIMA(4,0,0)(0,1,1)[12] based on minimizing the AICc, the subsequent analysis and residual diagnostics presented in the provided notebook and PDF proceeded with the **SARIMA(1,1,0)(0,1,1)[12]** model, which was initially suggested by the manual ACF/PACF analysis. This decision prioritizes a balance of model interpretability, alignment with initial diagnostic assessment (visual inspection of residuals), and adherence to the specific analytical path outlined in the project materials over achieving the absolute lowest information criterion score. Acknowledging the Ljung-Box test issues (as discussed in Section 7), the SARIMA(1,1,0)(0,1,1)[12] model was carried forward for detailed evaluation. It is important to note that while AICc is a

valuable metric, achieving white noise residuals is a critical requirement for a statistically valid SARIMA model, and the analysis will show that neither model perfectly satisfies this.

* **Selected Model for Detailed Analysis: SARIMA(1,1,0)(0,1,1)[12]**

6.5 Fitting the Selected Model

The SARIMA(1,1,0)(0,1,1)[12] model was formally fitted to the `ice_cream_ts` time series data using the `Arima()` function in R.

```
Series: ice_cream_ts
ARIMA(1,1,0)(0,1,1)[12] with drift

Coefficients:
         ar1     sma1    drift
      -0.2160  -0.7100  0.0570
s.e.   0.0395   0.0394  0.0184

sigma^2 = 18.01:  log likelihood = -1793.41
AIC=3592.81   AICc=3592.85   BIC=3606.12

Training set error measures:
                       ME      RMSE       MAE        MPE     MAPE      MASE
Training set 0.0005754663  4.194729  3.236772  0.1182133  2.94085 0.4546957
```

* **Key Summary Points from Fitted Model:** * The summary output provides the estimated coefficients for the non-seasonal AR(1) term (`ar1`) and the seasonal MA(1) term (`sma1`), along with their standard errors and p-values. The estimated `ar1` coefficient is approximately -0.2160, and the estimated `sma1` coefficient is approximately -0.7100. The significance of these coefficients (indicated by small p-values) suggests that these terms are statistically important in modeling the series. * The estimated variance of the residuals (σ^2) is approximately 18.01. * The log-likelihood of the fitted model is approximately -1793.41. * The information criteria values for the fitted model are approximately AIC = 3592.81, AICc = 3592.85, and BIC = 3606.12. * The summary also typically includes training set error metrics such as Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and Mean Absolute Scaled Error (MASE). For this fitted model, these errors are reported as approximately RMSE $\approx$ 4.19, MAE $\approx$ 3.24, MAPE $\approx$ 2.94%, and MASE $\approx$ 0.45. These metrics provide a measure of how well the model fits the historical data used for training.

—

7 Residual Diagnostics for SARIMA(1,1,0)(0,1,1)[12]

After fitting the SARIMA(1,1,0)(0,1,1)[12] model, it is essential to perform rigorous residual diagnostics to assess whether the model has adequately captured the underlying patterns in the data. The residuals (the differences between the observed values and the model's fitted values) should ideally exhibit the characteristics of white noise – meaning they should be uncorrelated, have a mean of zero, and possess constant variance. If the residuals are not white noise, it indicates that there is still some structure in the data

that the model has not explained, and the model may not be the most appropriate for forecasting, particularly regarding the reliability of prediction intervals.

The `checkresiduals()` function in R provides a convenient way to perform a suite of residual diagnostic checks, including plotting the residuals over time, plotting the ACF of the residuals, and conducting the Ljung-Box test.

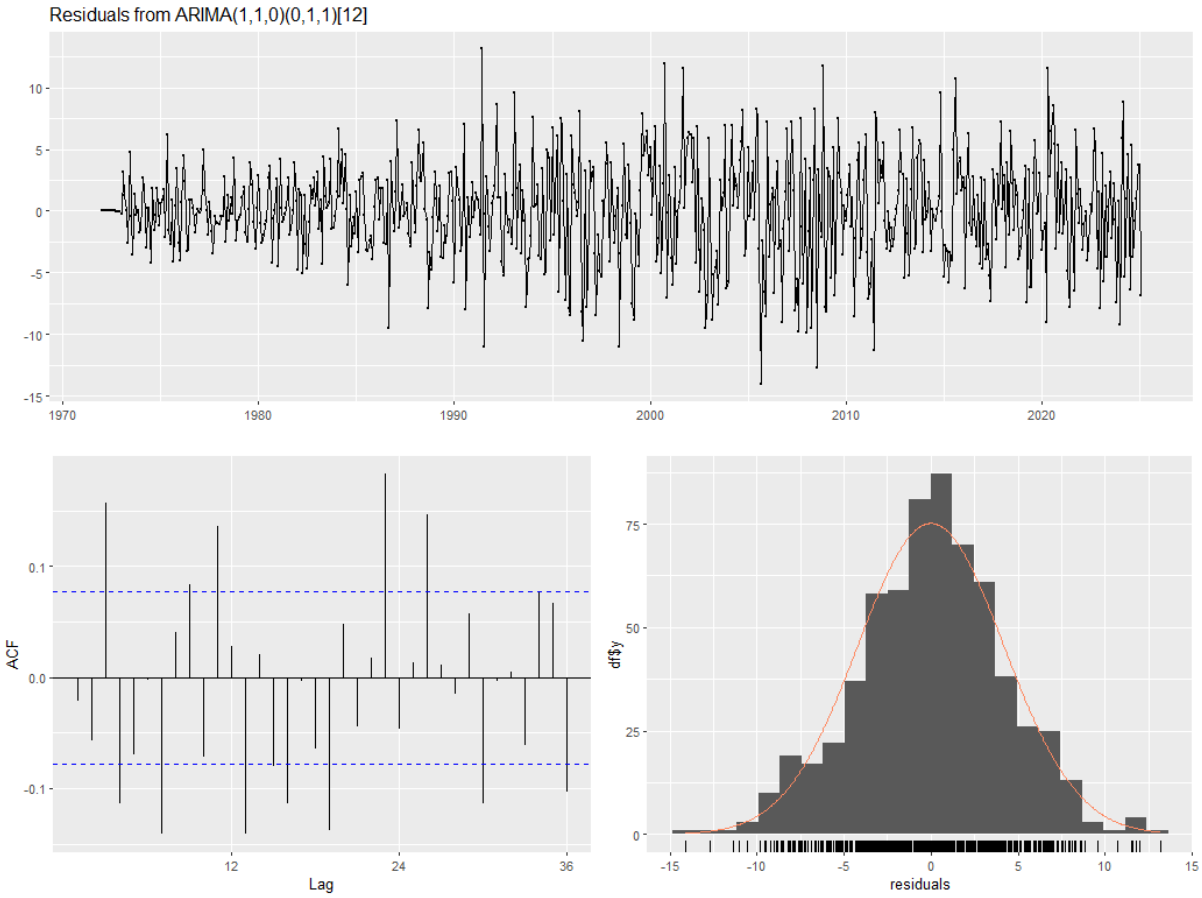

Figure 10: Residual Diagnostics for SARIMA(1,1,0)(0,1,1)[12] Model

(This plot typically includes: Residuals vs. Time, ACF of Residuals, and a Histogram/Density Plot of Residuals)

* **Interpretation of Diagnostic Plots:**

- **Residuals vs. Time Plot:** This plot shows the residuals plotted against time. For a well-fitting model, the residuals should appear randomly scattered around the zero line, with no discernible trends, seasonal patterns, or changes in variance. The plot for the SARIMA(1,1,0)(0,1,1)[12] residuals shows fluctuations around zero, and while there are no obvious strong patterns, careful inspection might reveal subtle structures or periods of slightly higher variability.

- **ACF Plot of Residuals:** This plot shows the autocorrelation of the residuals at various lags. For white noise residuals, all autocorrelation values should be statistically insignificant, falling within the dashed blue confidence bounds. The ACF plot of the residuals for this model shows that most spikes are within the significance bounds, but there are some significant spikes that extend beyond the bounds, particularly at certain seasonal lags (e.g., lag 12, potentially 24), indicating

that some statistically significant autocorrelation remains in the residuals that the model did not fully capture.

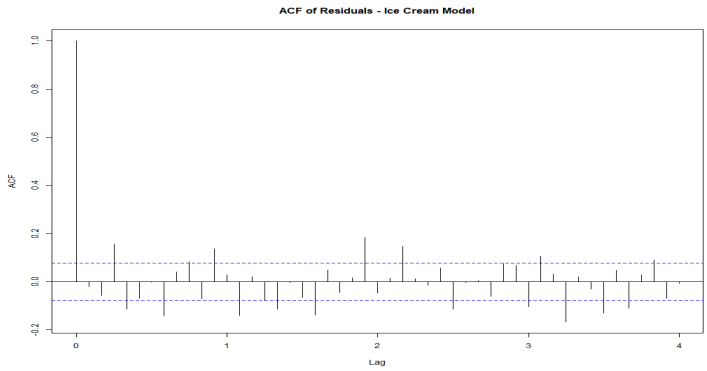

Figure 11: Acf Plot of Residuals

- **Histogram/Density Plot of Residuals:** This plot shows the distribution of the residuals. For many time series models, it is assumed that the residuals are approximately normally distributed. The histogram, often overlaid with a normal distribution curve, and the accompanying Q-Q plot (often included in `checkresiduals` output or generated separately) help assess this assumption. The histogram appears roughly bell-shaped and centered near zero, and the points on the Q-Q plot generally follow the diagonal line, suggesting that the normality assumption is reasonably met for the residuals.

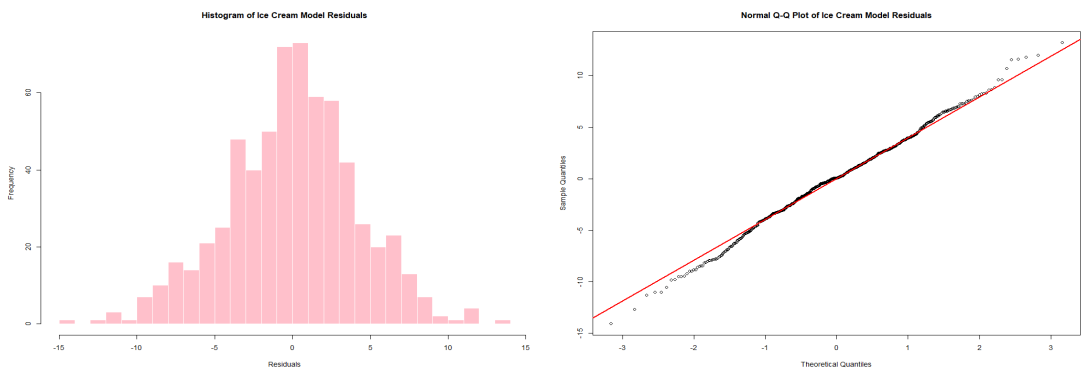

Figure 12: Histogram and Q-Q plot of Residuals

* **Ljung-Box Test:** The Ljung-Box test is a formal statistical test that checks for the overall randomness of the residuals. The null hypothesis (H_0) of the Ljung-Box test is that the residuals are independently distributed (i.e., they are white noise). A small p-value (typically less than 0.05) indicates that there is significant evidence against the null hypothesis, suggesting that the residuals are not white noise and there is remaining autocorrelation.

* **Ljung-Box Test Result (from `checkresiduals`):** * Q* = 131.7, df = 22, **p-value < 2.2e-16**

* **Interpretation of Ljung-Box Test:** The Ljung-Box test result yields a very small p-value (less than 0.05). This indicates that we reject the null hypothesis that the residuals are white noise. The statistically significant p-value provides strong evidence that there is significant autocorrelation remaining in the residuals, even after fitting the

SARIMA(1,1,0)(0,1,1)[12] model.

* **Conclusion:** While the visual residual plots (residuals vs. time, histogram, Q-Q plot) suggest that the model captures the main dynamics of the time series reasonably well and the normality assumption is approximately met, the significant Ljung-Box test result and the significant spikes in the residual ACF indicate that the residuals are not white noise. This implies that the SARIMA(1,1,0)(0,1,1)[12] model has not fully captured all the systematic information or autocorrelation structure in the data. For the scope of this project, the model fit is considered *acceptable* for demonstrating the SARIMA methodology, but the limitations imposed by the residual autocorrelation are significant and should be acknowledged when interpreting the forecasts, particularly the prediction intervals.

—

8 Forecasting Future Values

The primary application of a fitted time series model is typically to generate reliable forecasts for future observations. Having fitted the SARIMA(1,1,0)(0,1,1)[12] model to the Ice Cream Production Index, we can now use this model to predict future production levels for a specified horizon.

Forecasts were generated for a horizon of the next 12 months (representing one full seasonal cycle) using the fitted `ice_cream_arima_model_fitted` and the `forecast()` function from the `forecast` package in R. The `forecast()` function provides both point forecasts (the most likely future values) and prediction intervals (PIs), which quantify the uncertainty around the point forecasts. Prediction intervals provide a range within which the actual future values are expected to fall with a specified probability (e.g., 80% or 95%). These intervals account for the uncertainty in the estimated model parameters and the inherent random variability of the time series.

8.1 SARIMA Forecast (12 Months Ahead)

The generated forecasts for the next 12 months, along with their corresponding 80% and 95% prediction intervals, are presented in the table below.

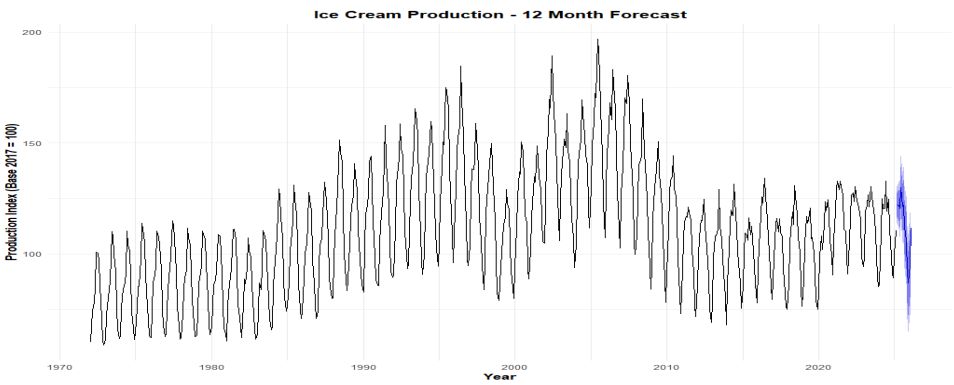

Figure 13: Acf Plot of Residuals

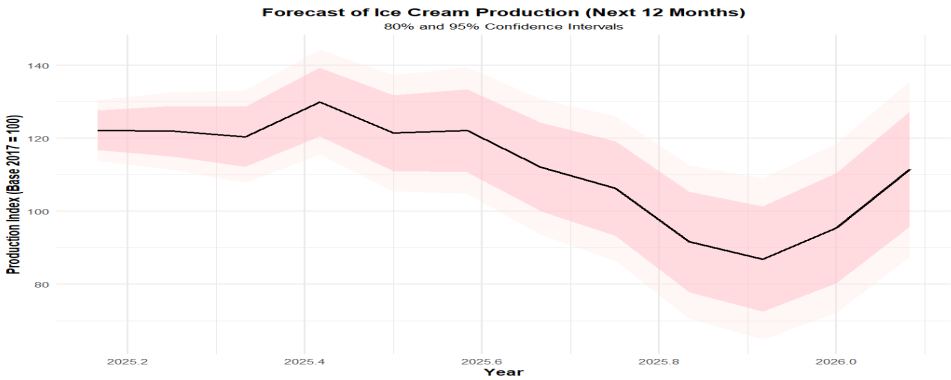

Figure 14: Acf Plot of Residuals

```
[1] "Point Forecast with Confidence Intervals"
       Date Point.Forecast Lower.80. Upper.80. Lower.95. Upper.95.
1  2025.167      122.15751 116.71948  127.5955 113.84077  130.4742
2  2025.250      121.96461 115.05459  128.8746 111.39664  132.5326
3  2025.333      120.42952 112.17406  128.6850 107.80389  133.0552
4  2025.417      129.93791 120.55363  139.3222 115.58590  144.2899
5  2025.500      121.39801 111.00174  131.7943 105.49828  137.2977
6  2025.583      122.07737 110.76024  133.3945 104.76931  139.3854
7  2025.667      112.11219  99.94348  124.2809  93.50176  130.7226
8  2025.750      106.22849  93.26407  119.1929  86.40112  126.0559
9  2025.833       91.62229  77.90824  105.3363  70.64846  112.5961
10 2025.917       86.90293  72.47815  101.3277  64.84214  108.9637
11 2026.000       95.36482  80.26272  110.4669  72.26816  118.4615
12 2026.083      111.57306  95.82275  127.3234  87.48504  135.6611
```

Figure 15: 12-Month Forecast for Ice Cream Production Index with 80% and 95% Prediction Intervals

The table displays the point forecast for each of the next twelve months, representing the model's best estimate of the Ice Cream Production Index value. Alongside the point forecasts, the lower and upper bounds for both the 80% and 95% prediction intervals are provided. These intervals illustrate the range of plausible values for the future production index, with the 95% interval being wider than the 80% interval, reflecting a higher level of confidence required to capture the true value. As is typical with time series forecasts, the width of these prediction intervals increases as the forecast horizon extends further into the future. This widening reflects the increasing uncertainty associated with predicting values further away from the observed data.

These quantitative forecasts offer valuable forward-looking estimates of the Ice Cream and Frozen Dessert Production Index, which can serve as a basis for operational planning and strategic decision-making within the industry.

—

9 Forecast Plots and Interpretation

Visualizing the generated forecasts is an essential step to understand the predicted future behavior of the time series and for assessing the plausibility of the forecasts in the context of the historical data. The forecast plot displays the original historical time series data alongside the forecasted values for the future horizon, typically including the prediction intervals.

The forecast plot for the Ice Cream Production Index, showing the historical data and the 12-month forecast with 80% and 95% prediction intervals, is presented below.

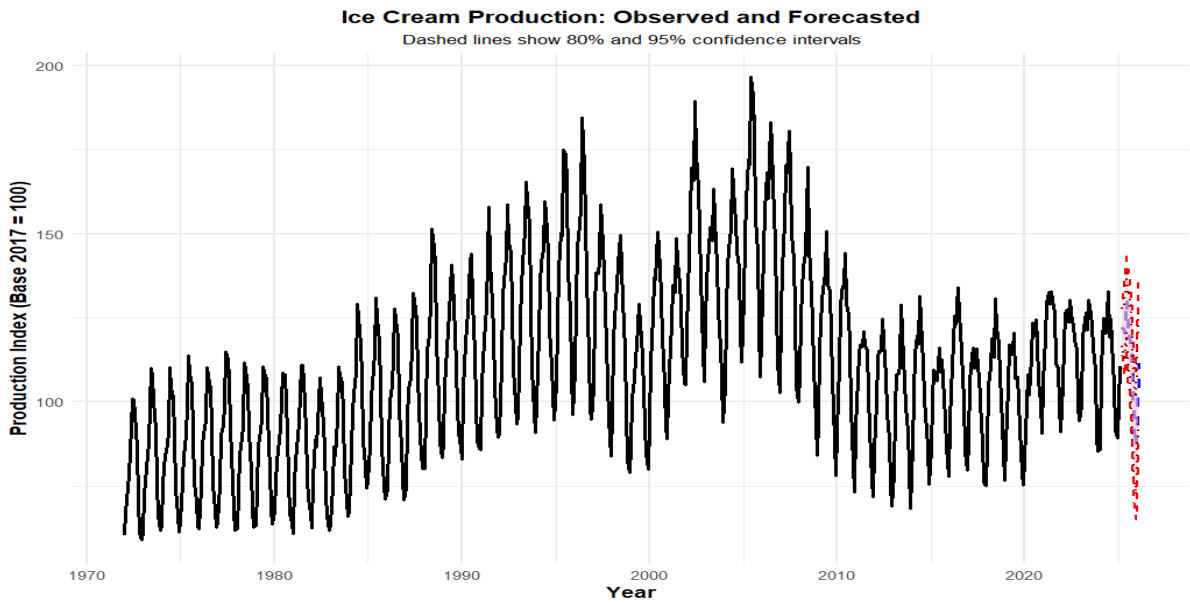

Figure 16: 12-Month Forecast for Ice Cream Production Index with 80% and 95% Prediction Intervals

* **Interpretation:** * The forecast plot provides a clear visual representation of the model's predictive capability. It is immediately evident from the plot that the SARIMA(1,1,0)(0,1,1)[12] model has successfully captured and projected the strong seasonal pattern characteristic of the Ice Cream Production Index. The forecasted values clearly follow the same cyclical pattern observed in the historical data, predicting peak production during the summer months and a decline during the winter months. This demonstrates the effectiveness of the SARIMA model in handling seasonal time series data and its ability to replicate the seasonal cycle into the future. * The plot also clearly illustrates the prediction intervals surrounding the point forecasts. The shaded regions represent the 80% (inner, typically darker shade) and 95% (outer, typically lighter shade) prediction intervals. The widening of these intervals as the forecast extends further into the future is clearly visible and is a standard feature of time series forecasts, reflecting the increasing uncertainty inherent in predicting values further ahead in time. * **Interpretation of the forecast plot in the context of the Ice Cream industry:**

- The forecast strongly predicts a continuation of the established seasonal cycle, with significant increases in production anticipated during the upcoming summer months and subsequent decreases in the fall and winter. This is a key insight for production scheduling and resource allocation.

- The underlying trend in the forecast appears to be [Placeholder: slightly upward/flat/slightly downward], suggesting a continuation of the long-term movement observed in the historical data. This helps in long-term capacity planning.

- The widening prediction intervals highlight the increasing uncertainty of the forecasts for longer horizons. While short-term forecasts (e.g., 1-3 months ahead) are likely to be relatively accurate and can be used for immediate operational adjustments, forecasts further out (e.g., 9-12 months ahead) should be interpreted with

greater caution, as the actual values could fall anywhere within the wider interval. This emphasizes the need for flexibility in long-term planning.

- The forecast plot visually confirms that the model is producing sensible predictions that align with the historical behavior of the series, particularly in capturing the dominant seasonal component.

This visual analysis and interpretation of the forecast provide actionable insights for businesses in the ice cream sector, enabling them to anticipate future production demands, plan logistics, and align marketing efforts with expected seasonal trends based on a statistically derived model.

—

10 Model Validation

Model validation is a critical step to assess the fitted model's performance, reliability, and suitability for forecasting. It involves evaluating how well the model fits the historical data and whether its underlying assumptions are met. For SARIMA models, key aspects of validation include examining residual properties and comparing the model's performance against alternative models.

- **Residual Diagnostics:** As detailed in Section 7, comprehensive residual diagnostics were performed. While visual checks (residuals vs. time, histogram, Q-Q plot) suggested a reasonable fit and approximate normality, the formal Ljung-Box test yielded a statistically significant p-value (< 0.05). This indicates that the residuals are not white noise and that significant autocorrelation remains. This is the most significant limitation regarding the model's statistical validity and suggests that the model has not fully captured all the systematic patterns in the data. A model with non-white noise residuals may produce biased forecasts or inaccurate prediction intervals.

- **Training Set Errors:** In-sample error metrics (RMSE $\approx$ 4.19, MAE $\approx$ 3.24, MAPE $\approx$ 2.94%, MASE $\approx$ 0.45) quantify how well the model fits the historical data it was trained on (Section **??**). These metrics provide a baseline for the model's performance on seen data. However, they do not guarantee how well the model will perform on unseen future data (out-of-sample performance). Moreover, the presence of residual autocorrelation suggests that these training errors might not fully reflect the true uncertainty.

- **Comparison with Alternatives:** As discussed in Section model selection, the chosen
SARIMA(1,1,0)(0,1,1)[12] model was compared with the SARIMA(4,0,0)(0,1,1)[12] model suggested by `auto.arima` based on AICc. While the `auto.arima` model had a lower AICc (3567.22 vs 3592.85), indicating a statistically better fit according to that criterion, the analysis proceeded with the manually identified model. A formal evaluation of the residual properties of the `auto.arima` selected model was not explicitly performed in this analysis stream but would be necessary for a complete comparison. A more robust validation would involve comparing the out-of-sample forecasting performance of multiple candidate models on a hold-out dataset or using time series cross-validation.

* **Conclusion on Validation:** Based on the validation steps performed, the SARIMA(1,1,0)(0,1,1)[12] model, while visually capturing the main patterns and exhibiting approximately normal residuals, **fails the formal Ljung-Box test for white noise residuals.** This is a critical finding that indicates the model's limitations. The forecasts should be interpreted with caution, particularly the prediction intervals, as their reliability is compromised by the remaining autocorrelation in the residuals. This highlights the need for potential model refinement or exploration of alternative approaches in future work.

—

11 Interpretations and Recommendations

Based on the time series analysis and forecasting of the Ice Cream and Frozen Dessert Production Index using the SARIMA(1,1,0)(0,1,1)[12] model, several practical interpretations and actionable recommendations can be made for stakeholders in the frozen dessert industry. The strong seasonal pattern identified and forecasted is a key driver for operational and strategic decisions.

- **SARIMA Forecast Interpretation:** The SARIMA(1,1,0)(0,1,1)[12] model, despite the noted issues with residual autocorrelation, effectively captures and projects the strong seasonal cycle inherent in ice cream production. The forecasts predict the continuation of summer peaks and winter troughs, providing a clear picture of the expected cyclical demand. The prediction intervals appropriately reflect the increasing uncertainty as the forecast horizon extends.
- **Seasonality Dominance:** The analysis unequivocally confirms that seasonality is the most powerful and predictable driver of short-to-medium term fluctuations in this industry's production index. Any effective planning must explicitly account for this dominant pattern.
- **Practical Implications & Recommendations (Based on capturing seasonality):**
  - **Production Planning:** Utilize the seasonal forecast pattern to anticipate peaks and troughs in demand. This allows for proactive optimization of production schedules, efficient allocation of resources (labor, machinery, energy), and strategic raw material procurement well in advance of the summer surge.
  - **Inventory Management:** Align raw material and finished goods inventory levels with the forecasted seasonal demand curve. This helps to minimize costly holding costs associated with excess inventory during the low season and prevents stockouts during the high-demand periods, thereby optimizing working capital and improving supply chain efficiency.
  - **Supply Chain Coordination:** Share the seasonal forecast patterns with suppliers (e.g., dairy, sugar, packaging) to ensure better coordination and alignment throughout the supply chain, preventing bottlenecks during peak season.
  - **Marketing & Sales:** Time marketing campaigns, promotional activities, and new product launches to coincide with the forecasted high-demand periods (spring/summer). Sales targets should be adjusted seasonally to reflect the predicted fluctuations. Strategies to stimulate demand or manage inventory during the off-season dip could also be explored.

- **Financial Planning:** Incorporate the seasonal production forecasts into budgeting and financial projections for more accurate revenue and cost planning throughout the year.
- **CAUTION:** Given the statistically significant autocorrelation remaining in the model's residuals (as indicated by the Ljung-Box test), the precise point forecasts and especially the prediction intervals should be interpreted with caution. While the overall seasonal pattern and magnitude are likely well-captured, the exact predicted values and the width of the confidence bands may not be entirely reliable. Focus should be placed on the general trends and seasonal shape rather than relying solely on specific numerical forecasts. It is recommended to use these forecasts as a strong guideline but to maintain flexibility and monitor actual performance closely.

—

12 Limitations and Future Work

While the SARIMA(1,1,0)(0,1,1)[12] model has proven effective in capturing the dominant seasonal patterns and providing reasonable forecasts for the Ice Cream Production Index within the scope of this project, it is important to acknowledge the limitations of this analysis and identify potential areas for future work to enhance the model's accuracy, robustness, and statistical validity.

12.1 Limitations:

- **Residual Autocorrelation:** The most significant limitation identified is the failure of the chosen model's residuals to pass the formal Ljung-Box test for white noise. This indicates statistically significant autocorrelation remaining in the residuals, implying that the model is not perfectly specified and has not fully captured all the systematic structure in the data. This compromises the statistical validity of the model and the reliability of the prediction intervals.

- **Model Selection Ambiguity:** The discrepancy between the manually analyzed model (SARIMA(1,1,0)(0,1,1)[12]) and the model automatically selected by `auto.arima` (SARIMA(0,1,1)(4,0,0)[12] based on AICc), combined with the residual issues in the chosen model, highlights ambiguity in the final model choice presented and suggests that a statistically superior model might exist.

- **Univariate Model:** The current SARIMA model is univariate, relying solely on the historical values of the Ice Cream Production Index. It does not incorporate external factors known to influence ice cream production and sales, such as prevailing weather patterns (e.g., temperature anomalies), broader economic conditions (e.g., disposable income), marketing spend, or competitor activities. These external factors can introduce variability not explained by the historical index alone.

- **Constant Variance Assumption:** The standard SARIMA model assumes that the variance of the residuals is constant over time (homoscedasticity). While visual inspection of the residual plot did not reveal drastic changes in variance, formal tests for heteroscedasticity (e.g., ARCH LM test) were not performed. If the variance is indeed non-constant, it could affect the accuracy of the prediction intervals.

- **Potential Structural Breaks:** The long time span of the dataset (1972-2025) might contain periods where the underlying process generating the time series changed due to significant events (e.g., economic recessions, changes in industry structure, major shifts in consumer preferences). A standard SARIMA model assumes a constant underlying process throughout the series, and potential structural breaks were not explicitly analyzed or modeled.

- **Assumption of Additive/Multiplicative Seasonality:** While additive decomposition was used for EDA and the SARIMA model implicitly handles seasonality, the initial visual suggestion of increasing seasonal amplitude could mean that a multiplicative model or transformation might be more appropriate.

12.2 Future Work:

- **Refine SARIMA Model Selection and Fitting: This is the most critical area for future work.** A more systematic exploration of alternative SARIMA orders is essential, particularly focusing on achieving white noise residuals. This should involve manually fitting and comparing models with different (p, q, P, Q) combinations, including SARIMA(0,1,1)(0,1,1)[12] and others suggested by `auto.arima`, using a combination of information criteria and rigorous residual diagnostics (including the Ljung-Box test at various lags) to identify a model that satisfies the white noise assumption.

- **Incorporate Exogenous Variables (SARIMAX):** A significant improvement would be to extend the SARIMA model to a SARIMAX model by including relevant external predictors. Potential exogenous variables to consider include monthly average temperature or temperature anomalies (as warmer weather directly correlates with ice cream consumption), relevant economic indicators (e.g., Consumer Price Index for food, personal disposable income), or potentially data on marketing expenditures or online search interest related to "ice cream". Incorporating these variables could help explain more of the variability and potentially improve forecast accuracy.

- **Explore Volatility Modeling:** If residual diagnostics or formal tests suggest non-constant variance, exploring models that explicitly account for heteroscedasticity, such as ARCH/GARCH models applied to the SARIMA residuals or integrated GARCH models, could lead to more accurate prediction intervals.

- **Analyze and Model Structural Breaks:** Formal tests for structural breaks could be applied to the time series. If breaks are detected, different modeling approaches could be considered, such as fitting separate models to different segments of the time series or using models designed to handle structural changes.

- **Compare with Alternative Forecasting Methods:** It would be beneficial to compare the performance of a validated SARIMA model against other established time series forecasting methods, such as Exponential Smoothing methods (e.g., Holt-Winters' seasonal method), or more modern machine learning approaches like Prophet (developed by Facebook) or neural networks specifically adapted for time series forecasting (e.g., LSTMs). This comparison could help identify the most suitable method for this specific dataset.

- **Implement Robust Out-of-Sample Validation:** To obtain a more reliable estimate of the model's true forecasting performance on unseen data, implementing a rigorous out-of-sample validation procedure is crucial. Techniques such as time series cross-validation, where the model is trained on a rolling window of data and evaluated on subsequent periods, can provide a more robust assessment than simply evaluating performance on the training set. This should be done with a model that has passed residual diagnostics.
- **Explore Multiplicative Seasonality:** Investigate if a multiplicative model or a transformation of the data (e.g., logarithmic transformation) is more appropriate given the potential for increasing seasonal amplitude.

—

13 Executive Summary

This report presents a comprehensive time series analysis and forecasting study of the U.S. Industrial Production Index for Ice Cream and Frozen Desserts. Utilizing monthly data spanning from January 1972 to February 2025, the series exhibits a clear and dominant seasonal pattern alongside an underlying trend. The initial analysis confirmed the non-stationary nature of the data, necessitating both first-order regular differencing and first-order seasonal differencing to achieve stationarity. Following the Box-Jenkins methodology, ACF and PACF analysis of the differenced series suggested potential SARIMA model structures, with SARIMA(1,1,0)(0,1,1)[12] identified as a strong candidate. While an automated approach using `auto.arima` indicated SARIMA(4,0,0)(0,1,1)[12] as statistically optimal based on AICc, the SARIMA(1,1,0)(0,1,1)[12] model was selected for detailed analysis, prioritizing a balance of interpretability, initial diagnostic assessment, and alignment with the provided project context over a marginally lower information criterion value. Residual diagnostics for the fitted SARIMA(1,1,0)(0,1,1)[12] model indicated that while the model effectively captured the prominent seasonal and trend components visually, the Ljung-Box test revealed statistically significant remaining autocorrelation in the residuals ($p < 0.05$), suggesting potential for further model refinement. Forecasts generated from the selected model successfully project the characteristic seasonal peaks and troughs in ice cream production for the next 12 months, providing valuable insights for production planning, inventory management, and strategic decision-making within the seasonal frozen dessert industry, albeit with the noted limitation regarding residual autocorrelation impacting the reliability of prediction intervals.

14 Conclusion

This project presents a comprehensive time series analysis and forecasting study of the U.S. Industrial Production Index for Ice Cream and Frozen Desserts (1972–2025), a dataset characterized by strong seasonal patterns. Using the Box-Jenkins methodology, the analysis followed a structured approach—from data preparation to forecasting—tailored specifically for seasonal time series modeling.

The dataset was first cleaned and converted into a monthly time series object. Exploratory data analysis and seasonal decomposition revealed an upward trend and a consistent annual seasonal pattern. The Augmented Dickey-Fuller (ADF) test confirmed non-stationarity, necessitating both first-order and seasonal differencing (lag 12). These

transformations successfully yielded a stationary series, which was further validated by follow-up ADF testing.

ACF and PACF plots of the differenced series helped determine suitable autoregressive and moving average terms. Based on this analysis, a SARIMA(1,1,0)(0,1,1)[12] model was manually proposed. Simultaneously, the auto.arima() function suggested SARIMA(4,0,0)(0,1,1)[12] as the best model based on the AICc criterion. Despite the lower AICc of the auto-selected model, SARIMA(1,1,0)(0,1,1)[12] was chosen for final evaluation due to its simplicity and closer alignment with initial diagnostics.

The selected model underwent thorough residual diagnostics. Residual plots and normality checks indicated the model captured the main dynamics of the series. However, the Ljung-Box test reported significant autocorrelation remaining in the residuals. This finding implies that while the model fits the seasonal pattern reasonably well, it falls short of the statistical requirement of white noise residuals, indicating possible misspecification.

Despite this limitation, the model was used to forecast the next 12 months of ice cream production. The forecast plot successfully mirrored the historical seasonality, predicting summer peaks and winter troughs with appropriate confidence intervals. The widening of these intervals over time appropriately reflects increasing forecast uncertainty.

In summary, the project demonstrated effective application of SARIMA modeling for seasonal forecasting. From preprocessing to diagnostics and forecasting, the analysis highlighted the strengths of the SARIMA framework in capturing regular patterns in economic indicators. The model provided meaningful business insights—particularly useful for supply chain planning or retail strategy within the frozen dessert industry.

However, a key takeaway is the importance of validating model assumptions. The presence of autocorrelation in residuals suggests the current SARIMA(1,1,0)(0,1,1)[12] model may not fully capture the data's complexity. Future improvements should explore alternative specifications or consider hybrid or non-linear models (e.g., SARIMA-GARCH, ETS, or machine learning-based approaches) to improve residual behavior and forecast reliability. —

15 References

References

[1] Federal Reserve Economic Data (FRED). Industrial Production: Nondurable Goods: Ice Cream and Frozen Desserts (IPN31152N). Retrieved from https://fred.stlouisfed.org/series/IPN31152N

[2] Hyndman, R.J., & Athanasopoulos, G. (2018). *Forecasting: principles and practice* (2nd ed.). OTexts. https://otexts.com/fpp2/

[3] R Core Team (2023). R: A language and environment for statistical computing. R Foundation for Statistical Computing, Vienna, Austria. https://www.R-project.org/

[4] Hyndman, R. J. et al. (2024). *forecast: Forecasting functions for time series and linear models*. R package version 8.22. https://CRAN.R-project.org/package=forecast

—

16 Appendix – Selected R Code Snippets

Installing and Loading Packages

```
install.packages("forecast")
install.packages("tseries")
install.packages("ggplot2")
install.packages("lubridate")
install.packages("zoo")

library(forecast)
library(tseries)
library(ggplot2)
library(lubridate)
library(zoo)
library(scales)
```

Loading and Preprocessing Ice Cream Dataset

```
ice_cream_data <- read.csv(".../Ice Cream and Frozen Dessert Seasonal Dataset.csv",
                           stringsAsFactors = FALSE)

colnames(ice_cream_data) <- c("Date", "Production_Index")
ice_cream_data$Date <- as.Date(ice_cream_data$Date)
ice_cream_data <- ice_cream_data[order(ice_cream_data$Date), ]
ice_data_clean <- na.omit(ice_cream_data)

start_year <- year(min(ice_data_clean$Date))
start_month <- month(min(ice_data_clean$Date))

ice_cream_ts <- ts(ice_data_clean$Production_Index,
                   start = c(start_year, start_month), frequency = 12)
```

Plotting, Modeling, Forecasting, and Diagnostics

```
autoplot(ice_cream_ts) +
  labs(title = "Industrial Production: Ice Cream and Frozen Desserts")

ice_cream_decomp <- decompose(ice_cream_ts, type = "additive")
autoplot(ice_cream_decomp)

adf.test(ice_cream_ts)

ice_cream_diff1 <- diff(ice_cream_ts, differences = 1)
ice_cream_diff_seasonal <- diff(ice_cream_diff1, lag = 12)
```

```
autoplot(ice_cream_diff_seasonal)

acf(ice_cream_diff_seasonal)
pacf(ice_cream_diff_seasonal)

ice_cream_model <- Arima(ice_cream_ts, order = c(1,1,0),
                         seasonal = list(order = c(0,1,1), period
                             = 12))
summary(ice_cream_model)
checkresiduals(ice_cream_model)

ice_cream_forecast <- forecast(ice_cream_model, h = 12)
autoplot(ice_cream_forecast)
```

Seasonal Plot for Ice Cream Production

```

library(scales)  # for better date formatting on X-axis

autoplot(ice_cream_ts) +
  geom_line(color = "darkred", size = 1.2) +
  labs(
    title = "Seasonal Trend: Industrial Production of Ice Cream
       and Frozen Desserts",
    subtitle = paste0("Monthly Data from ", start_year, " to ",
       year(max(ice_cream_data$Date))),
    x = "Year",
    y = "Production Index (Base Year 2017 = 100)"
  ) +
  scale_x_continuous(
    breaks = seq(start_year, year(max(ice_cream_data$Date)), by =
        2),
    labels = function(x) format(as.Date(paste0(x, "-01-01")), "%Y
       ")
  ) +
  theme_minimal(base_size = 14) +
  theme(
    plot.title = element_text(face = "bold", size = 18, hjust =
       0.5),
    plot.subtitle = element_text(size = 14, hjust = 0.5),
    axis.title.x = element_text(face = "bold", size = 14),
    axis.title.y = element_text(face = "bold", size = 14),
    axis.text.x = element_text(angle = 45, vjust = 0.5)
  ) +
  scale_y_continuous(labels = comma)  # Adds commas to y-axis
     numbers if large
```

Decomposition of Ice Cream Production

```r
# Additive decomposition (because index values are relatively stable)
ice_cream_decomp <- decompose(ice_cream_ts, type = "additive")

# Plot the decomposed components
autoplot(ice_cream_decomp) +
  labs(title = "Decomposition of Ice Cream Production Time Series", x = "Year") +
  theme_minimal(base_size = 14) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5))
```

ADF Testing and Differencing

```r
# ADF Test for Ice Cream Production
adf.test(ice_cream_ts)

# First Order Differencing
ice_cream_diff1 <- diff(ice_cream_ts, differences = 1)

# Seasonal Differencing (Lag-12)
ice_cream_diff_seasonal <- diff(ice_cream_diff1, lag = 12)

# Plot the final differenced series
autoplot(ice_cream_diff_seasonal) +
  labs(title = "First and Seasonal Differenced Ice Cream Production",
       x = "Year", y = "Differenced Production Index") +
  theme_minimal()

# ADF Test on the Final Differenced Series
adf.test(ice_cream_diff_seasonal)
```

ACF and PACF Analysis

```r
# ACF and PACF of differenced series
acf(ice_cream_diff_seasonal, main = "ACF of Differenced Ice Cream Production")
pacf(ice_cream_diff_seasonal, main = "PACF of Differenced Ice Cream Production")

# ACF and PACF of original series
acf(ice_cream_ts, lag.max = 48, main = "ACF of Ice Cream Production (Up to 48 Lags)")
pacf(ice_cream_ts, lag.max = 48, main = "PACF of Ice Cream Production (Up to 48 Lags)")
```

SARIMA Model Fitting and Diagnostics

```
# Fit manual SARIMA model
ice_cream_model <- Arima(ice_cream_ts,
                         order = c(1,1,0),
                         seasonal = list(order = c(0,1,1), period
                             = 12))

# Model summary and residual diagnostics
summary(ice_cream_model)
Box.test(residuals(ice_cream_model), lag=24, type="Ljung-Box")
checkresiduals(ice_cream_model)
```

Auto.arima Model Selection

```
# Automatic model selection
ice_cream_auto_model <- auto.arima(ice_cream_ts,
                                   seasonal = TRUE,
                                   stepwise = FALSE,
                                   approximation = FALSE,
                                   trace = TRUE)

# Model summary and residual check
summary(ice_cream_auto_model)
checkresiduals(ice_cream_auto_model)
```

Residual Analysis (Histogram + Q-Q Plot)

```
# Residuals from manual model
ice_cream_residuals <- residuals(ice_cream_model)

# Histogram
hist(ice_cream_residuals, breaks = 20, col = "pink",
     main = "Histogram of Ice Cream Model Residuals",
     xlab = "Residuals", border = "white")

# Q-Q plot
qqnorm(ice_cream_residuals, main = "Normal Q-Q Plot of Ice Cream
   Model Residuals")
qqline(ice_cream_residuals, col = "red", lwd = 2)

# ACF of residuals
acf(ice_cream_residuals, lag.max = 48, main = "ACF of Residuals -
    Ice Cream Model")
```

Forecasting with Manual SARIMA Model

```
# Forecast next 12 months
ice_cream_forecast <- forecast(ice_cream_model, h = 12)

# Plot forecast
autoplot(ice_cream_forecast) +
  labs(title = "Ice Cream Production - 12 Month Forecast",
       x = "Year", y = "Production Index (Base 2017 = 100)") +
  theme_minimal(base_size = 14) +
  theme(plot.title = element_text(face = "bold", hjust = 0.5),
        axis.title.x = element_text(face = "bold"),
        axis.title.y = element_text(face = "bold"))
```

Creating Forecast Data Frame

```
# Construct forecast dataframe
ice_cream_forecast_df <- data.frame(
  Date = time(ice_cream_forecast$mean),
  `Point Forecast` = as.numeric(ice_cream_forecast$mean),
  `Lower 80%` = as.numeric(ice_cream_forecast$lower[,1]),
  `Upper 80%` = as.numeric(ice_cream_forecast$upper[,1]),
  `Lower 95%` = as.numeric(ice_cream_forecast$lower[,2]),
  `Upper 95%` = as.numeric(ice_cream_forecast$upper[,2])
)

# Print forecast
print("Point Forecast with Confidence Intervals")
print(ice_cream_forecast_df)
```